

When Blocks Build Societies
Multi-Agent Systems in Minecraft

Jan-Philipp Kiel
University of Cologne

Advanced Seminar
Information Systems and Digital Technology
Prof. Dr. Stefan Seidel

July 29, 2025

Abstract

This paper presents a comparative analysis of multi-agent systems (MAS) within the Minecraft environment. Focusing on Project Sid and TeamCraft as case studies, the research investigates how task-oriented and socially-oriented paradigms address the challenge of sustainable collaboration among autonomous agents. Project Sid exemplifies emergent, self-organizing agent societies, where roles, norms, and communication patterns develop organically, enabling adaptability and long-term sustainability. In contrast, TeamCraft employs structured, task-driven coordination with explicit role assignments and protocols, optimizing efficiency for well-defined objectives but limiting flexibility in novel contexts. The study employs a structured analytical framework to compare organizational structures, communication and coordination mechanisms, learning and adaptation strategies, and scalability. Findings highlight fundamental trade-offs: emergent approaches foster resilience and adaptability but may sacrifice initial efficiency, while structured frameworks excel in familiar scenarios but struggle with generalization and long-term viability. The analysis yields design principles for robust MAS, emphasizing the integration of both structured and emergent elements, effective communication, and ethical oversight. The insights extend beyond Minecraft, informing the development of adaptable, sustainable, and ethically aligned MAS in diverse virtual and real-world domains.

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

Table of Contents

Abstract	I
Table of Contents	II
Table of Figures	III
Table of Tables	IV
1 Introduction	1
2 Background and Related Work.....	2
2.1 Multi-Agent Systems in Virtual Environments	2
2.2 Minecraft as a Research Platform.....	3
2.3 Approaches to Agent Collaboration.....	4
3 Methodology	5
3.1 Case Study Selection.....	6
3.2 Literature Search and Screening	6
3.3 Analytical Framework	7
3.4 Use of Generative AI	8
4 Case Studies	9
4.1 Project Sid	9
4.2 TeamCraft	10
5 Comparative Analysis	11
5.1 Organizational Structures	11
5.2 Communication and Coordination.....	12
5.3 Learning and Adaptation.....	13
5.4 Sustainability and Scalability	14
6 Discussion.....	15
7 Conclusion and Future Work	16
References	17
A Figures and Tables	20
B AI Conversation Logs	27
Declaration of Generative AI Tool Usage.....	34

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

Table of Figures

Figure 1: An example Minecraft technology dependency tree for the mining of gold, diamond, and emeralds (AL et al., 2024).	20
Figure 2: Illustration of trading activities among agents in a Minecraft multi-agent system (AL et al., 2024).	20
Figure 3: Visual abstract of Project Sid, illustrating the progression from agent architecture to emergent agent civilization (AL et al., 2024).	21
Figure 4: Example of a multi-modal prompt for building in Minecraft (Long et al., 2024).	22

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

Table of Tables

Table 1: Comparison of candidate case studies with features and selection rationale.....	23
Table 2: Comparison of Project Sid and TeamCraft across key dimensions.	25

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

1 Introduction

Multi-agent systems (MAS) are a key frontier in artificial intelligence (AI) research, offering insights into how autonomous entities can collaborate, compete, and coexist within shared environments (Shu et al., 2024; van Leeuwen, 2021). These systems enable the study of complex emergent behaviors and collective intelligence that extend beyond the capabilities of single agents. As AI continues to advance, understanding the principles that govern effective multi-agent collaboration becomes increasingly vital for applications ranging from robotic teams to virtual assistants.

Despite advances in MAS, there is limited understanding of how different collaborative frameworks—task-oriented versus socially-oriented—impact sustainable collaboration in complex virtual environments such as Minecraft. This paper addresses the need to compare these approaches to inform future design of collaborative MAS. The scope of this research is limited to Minecraft-based MAS. While potential parallels to other domains are discussed, all findings and design principles are grounded in the Minecraft context.

Minecraft, with its open-world sandbox, procedurally generated environments, and flexible crafting, make it an ideal platform for multi-agent research (Fan et al., 2022). Its block-based world allows for precise control while still presenting meaningful challenges in navigation, resource gathering, construction, and social interaction (Zhao et al., 2024).

Most MAS research in Minecraft has focused on specific task completion, such as navigation or building predefined structures. However, a significant gap exists between these narrowly defined task-oriented systems and more complex, socially-oriented agent societies (Qiu et al., 2024; Shu et al., 2024). Bridging this gap requires new approaches to agent architecture, communication protocols, and coordination mechanisms that can support sustainable collaboration over extended periods.

Two recent research initiatives have made significant progress in addressing different aspects of this challenge: Project Sid and TeamCraft. Project Sid, introduced in 2024, explores large-scale simulations of 10-1000+ AI agents, focusing on how these agents develop specialized roles, adhere to collective rules, and engage in cultural transmission within a civilization-inspired framework (AL et al., 2024). In contrast, TeamCraft provides a multi-modal multi-agent benchmark featuring 55,000

task variants designed to test various aspects of agent collaboration through structured prompts and demonstrations (Long et al., 2024).

This paper addresses the research question: *How do different collaborative frameworks in Minecraft address the challenge of sustainable multi-agent collaboration within virtual environments?* By comparing these approaches, we aim to identify design principles that can inform the development of more robust and adaptable collaborative AI systems within virtual environments.

The paper is structured as follows: First, we review background literature on MAS, Minecraft as a research platform, and approaches to agent collaboration. Next, we detail our methodology, including case study selection criteria and analytical framework. We then analyze Project Sid and TeamCraft individually before conducting a comparative analysis across organizational structures, communication patterns, learning mechanisms, and sustainability. Finally, we discuss theoretical implications, explore practical applications in other domains, and conclude with directions for future research.

2 Background and Related Work

To ensure conceptual clarity, this paper adopts the following definitions: An *agent* is an entity perceiving and acting on its environment (Russell & Norvig, 2009); *artificial intelligence (AI)* is the study of agents maximizing cumulative reward (Russell & Norvig, 2009); a *multi-agent system (MAS)* consists of multiple interacting agents whose collective behavior emerges from their interactions (Wooldridge, 2009); and *collaboration* refers to agents working together to achieve shared goals, requiring coordination and communication (Panait & Luke, 2005).

2.1 Multi-Agent Systems in Virtual Environments

MAS represent a paradigm in AI research where multiple autonomous agents interact within a shared environment to achieve individual or collective goals (Shu et al., 2024; van Leeuwen, 2021). The field has evolved significantly from early distributed AI research in the 1980s to contemporary approaches leveraging deep learning and reinforcement learning. Modern MAS research encompasses diverse architectures ranging from fully cooperative to competitive and mixed-motive scenarios, with applications spanning robotics, logistics, gaming, and social simulation.

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

Virtual environments offer particularly valuable platforms for multi-agent research due to their controlled conditions, cost-effectiveness, and ability to simulate complex scenarios that would be impractical or impossible in physical settings (Lhakshmana et al., 2018; Shu et al., 2024). These environments provide a “consequence-free” space to test agent behaviors that might be risky or expensive in real-world settings, while still capturing essential dynamics of multi-agent interaction. Additionally, virtual environments allow researchers to precisely control experimental conditions and collect comprehensive data about agent behaviors and interactions.

Multi-agent architectures can be broadly categorized as centralized or decentralized, with various hybrid approaches emerging to leverage the advantages of both (van Leeuwen, 2021). Centralized approaches maintain a global view of the environment and coordinate agent actions through a central controller, offering optimal decision-making at the cost of scalability and robustness. Decentralized approaches distribute decision-making among individual agents, providing resilience and scalability at the expense of potentially suboptimal coordination. Communication models within these architectures range from direct message passing to stigmergic coordination through environmental modifications.

Recent advances in multi-agent collaboration have been driven by developments in large language models (LLMs) and deep reinforcement learning (Qiu et al., 2024; Shu et al., 2024). LLMs have enhanced agents’ ability to understand complex instructions and generate contextually appropriate responses, while reinforcement learning techniques have improved adaptation to dynamic environments and coordination among heterogeneous agents (Chen et al., 2024). These advances have expanded the scope of multi-agent research to include emergent communication protocols, collective intelligence, and social learning.

2.2 Minecraft as a Research Platform

Minecraft offers a uniquely suitable platform for AI research due to its procedurally generated worlds, comprehensive crafting system (see Figure 1), and physics simulation that strikes a balance between realism and abstraction (Fan et al., 2022). The game’s block-based structure provides a discretized state space that is more computationally tractable than continuous environments while still offering sufficient complexity to model interesting real-world problems. Furthermore, its open-world nature

allows for a wide range of tasks, from simple navigation to complex collaborative construction projects.

Previous AI research in Minecraft has explored diverse challenges, including single-agent task completion, navigation in procedurally generated landscapes, and creative building (Wang et al., 2025; Zhao et al., 2024). The MineDojo framework, for instance, provides thousands of diverse open-ended tasks and an internet-scale knowledge base with Minecraft videos, tutorials, and forum discussions to facilitate agent learning (Fan et al., 2022). Similarly, JARVIS-1 demonstrates how multimodal language models can perceive visual observations and textual instructions to generate sophisticated plans and perform embodied control in Minecraft (Wang et al., 2025).

The Minecraft environment presents unique challenges for AI research, including a large state space, partial observability, and the need for long-horizon planning (Fan et al., 2022; Wang et al., 2025). Agents must reason about objects that may be out of view, plan sequences of actions that may take hundreds or thousands of steps to complete, and adapt to unexpected environmental changes. These challenges make Minecraft particularly valuable for testing the limits of current AI approaches and driving innovations in planning, perception, and learning algorithms.

Beyond technical challenges, Minecraft enables social interaction and emergent behaviors that are difficult to study in more constrained environments (see Figure 2) (AL et al., 2024). The MineLand simulator, for example, introduces features like limited multimodal senses and physical needs to create more ecologically valid social interactions among large groups of agents (Yu et al., 2024). This capability for modeling complex social dynamics makes Minecraft an ideal test bed for studying how artificial societies might form, evolve, and sustain themselves over time.

2.3 Approaches to Agent Collaboration

Agent collaboration approaches typically fall into two broad categories: task-oriented collaboration, which focuses on completing specific objectives, and social-oriented collaboration, which emphasizes emergent social structures and norms (Qiu et al., 2024; Shu et al., 2024). Task-oriented approaches typically define explicit goals, roles, and protocols for interaction, prioritizing efficiency and measurable outcomes. Social-oriented approaches, in contrast, allow organizational structures and behavioral norms to emerge through repeated interactions, often leading to more adaptable but

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

less predictable collaborative behaviors.

Coordination mechanisms in MAS include role assignment, task allocation, and conflict resolution strategies (Dong et al., 2024; Shu et al., 2024). Role-based approaches assign specific responsibilities to agents based on their capabilities or position within an organizational structure. Task allocation methods distribute work among agents to optimize resource utilization and minimize conflicts. These mechanisms may be static (predetermined) or dynamic (adapting to changing conditions), with varying degrees of centralization in decision-making authority.

Communication strategies play a crucial role in effective collaboration, ranging from direct communication through message passing to indirect communication through environmental modifications (Qiu et al., 2024; Shu et al., 2024). Direct communication allows for precise information sharing but requires protocols for when, what, and how to communicate. Indirect communication (stigmergy) relies on agents observing and responding to changes in the shared environment, often leading to emergent coordination patterns. The choice of communication strategy significantly impacts a system’s scalability, robustness, and adaptability to novel situations.

Evaluation metrics for successful collaboration include task completion rates, resource efficiency, adaptability to changing conditions, and robustness to agent failures (Chen et al., 2024; Qiu et al., 2024; Shu et al., 2024). Task-oriented systems typically prioritize metrics related to goal achievement and efficiency, such as time-to-completion and resource consumption. Social-oriented systems may emphasize metrics related to sustainability and resilience, such as the stability of social structures and adaptability to unexpected challenges. These different evaluation priorities reflect fundamental differences in how collaboration is conceptualized across approaches.

3 Methodology

This study uses a comparative case study approach to analyze multi-agent collaboration in Minecraft, following established guidelines from Yin (2018) and Eisenhardt (1989). The process included: (1) identifying and screening candidate systems using explicit inclusion and exclusion criteria; (2) conducting a literature search across major academic databases; (3) applying a structured analytical framework to compare selected cases; and (4) documenting the use of generative AI tools, including ethical considerations and verification procedures. This ensures transparency, replicability,

and systematic analysis (Eisenhardt, 1989; Yin, 2018).

Project Sid and TeamCraft were selected as recent, well-documented, and conceptually contrasting case studies. Other systems were excluded due to overlap or insufficient distinctiveness. The analytical framework compares organizational structure, communication and coordination, learning and adaptation, and sustainability and scalability.

Generative AI tools supported literature review, synthesis, and drafting. All AI-generated content was reviewed and verified by the author. Limitations include potential selection bias, incomplete literature coverage, and risks of AI-generated inaccuracies. These were mitigated by explicit criteria, multi-database searches, and systematic verification against original sources.

3.1 Case Study Selection

Project Sid and TeamCraft were selected as the most methodologically robust and conceptually contrasting case studies of multi-agent collaboration in Minecraft. Selection followed criteria outlined by Yin (2018) and Eisenhardt (1989) and was based on recency, documentation quality, and clear contrast in collaborative paradigms. Other case studies were excluded for reasons such as overlap in paradigm, lack of focus on multi-agent collaboration, or insufficient distinctiveness. A detailed overview of all candidate case studies and selection rationale is provided in Table 1.

3.2 Literature Search and Screening

The literature search for candidate case studies was conducted between June and July 2025 using major academic databases including Google Scholar, Semantic Scholar, and arXiv. Search keywords included combinations of “multi-agent systems”, “Minecraft”, “collaboration”, “benchmark”, and “emergent behavior”. Initial results were screened for relevance based on title and abstract. Full texts were reviewed for systems that focused on multi-agent collaboration in Minecraft environments. Generative AI tools were used to assist in the screening and review of identified papers, enabling efficient processing and synthesis of the literature. Additionally, conversational learning with AI was used to quickly understand key concepts and frameworks; the conversation histories are referenced in the appendix (see Appendix B).

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

3.3 Analytical Framework

The comparative analysis employs a structured framework examining four key dimensions: (1) Organizational structure, (2) Communication and coordination, (3) Learning and adaptation, and (4) Sustainability and scalability. Each dimension is analyzed through the lens of established MAS literature and organizational theory, providing a systematic basis for comparison (Chen et al., 2024; Lhaksmana et al., 2018).

Organizational structure is relevant because the way agents are organized—whether through emergent hierarchies or predefined roles—directly influences the system’s ability to support sustainable collaboration and adapt to changing environments. Communication and coordination are central to effective multi-agent collaboration, as the mechanisms for information exchange and joint action determine how well agents can work together to achieve shared goals. Learning and adaptation are crucial for long-term sustainability, enabling agents to improve their behaviors and strategies in response to new challenges and evolving contexts. Sustainability and scalability address the research question by evaluating whether collaborative frameworks can maintain performance and resilience as the number of agents and complexity of tasks increase, which is essential for robust multi-agent societies in virtual environments like Minecraft.

Organizational structure is analyzed through the lens of hierarchy (the levels of authority and decision-making in the system), role specialization (how agents develop specialized functions), and adaptability (how the organization responds to changing conditions) (Chen et al., 2024; Lhaksmana et al., 2018). Project Sid’s emergent hierarchies are contrasted with TeamCraft’s task-oriented organization to understand how different structural approaches affect collaborative outcomes. This analysis draws on organizational theory concepts such as centralization, formalization, and departmentalization to characterize the systems’ organizational features .

Agent autonomy and communication patterns are evaluated by examining communication frequency (how often agents exchange information), content (what information is shared), and effectiveness (how communication impacts task performance) (Qiu et al., 2024; Shu et al., 2024). The analysis considers both direct communication through message passing and indirect communication through environmental modifi-

cations. Particular attention is paid to how communication patterns support or hinder coordination in different contexts, such as novel tasks or unexpected environmental changes.

Task allocation and sustainability are measured through metrics including efficiency (resource usage relative to task completion), adaptability (performance on novel tasks or with novel team compositions), and long-term viability (stability of performance over extended periods) (Chen et al., 2024; Shu et al., 2024). These metrics provide insight into how well each system balances short-term performance goals with long-term sustainability concerns, a critical consideration for enduring multi-agent collaborations.

3.4 Use of Generative AI

This research employed generative AI tools throughout the research process, including literature review, analysis, and manuscript preparation. Specifically, LLMs were used to assist in identifying relevant literature, synthesizing findings across sources, and generating initial drafts of paper sections. This approach allowed for efficient processing of the extensive literature on MAS while maintaining focus on the specific research question.

The primary AI tool used was a LLM with research capabilities, which provided assistance in summarizing academic papers, identifying connections between concepts, and generating coherent prose based on specified content requirements. The model was provided with search results from academic databases and asked to summarize key findings, compare approaches, and identify research gaps. Human oversight was maintained throughout this process, with all AI-generated content reviewed, revised, and expanded based on direct examination of the original sources.

Ethical considerations in the use of generative AI included ensuring proper attribution of ideas to original sources, verifying factual claims against primary literature, and maintaining transparency about the role of AI in the research process. All content generated by AI was treated as a first draft requiring human verification and refinement rather than as finished work. This approach aligns with emerging best practices for AI-assisted academic writing, which emphasize human oversight and accountability (Lund et al., 2023).

The use of generative AI introduces certain limitations and potential biases,

including the risk of propagating inaccuracies present in the training data, generating plausible-sounding but incorrect information, and prioritizing certain types of sources over others. To mitigate these risks, a systematic verification process was implemented, comparing AI-generated content against original sources and seeking additional sources where information appeared incomplete or potentially biased. The complete record of interactions with the AI system is included in Appendix B to ensure full transparency about the research process. All sources were properly cited, and all AI-generated content was reviewed and verified by the author to ensure accuracy and compliance with ethical standards for academic research and responsible AI use.

4 Case Studies

4.1 Project Sid

Project Sid is a large-scale simulation of 10 to over 1000 AI agents in Minecraft, designed to investigate the emergence of complex social structures and behaviors through extended interaction (see Figure 3) (AL et al., 2024). The system is built on the PIANO (Parallel Information Aggregation via Neural Orchestration) architecture, which enables agents to process multimodal inputs, maintain internal states, and interact in real time with both humans and other agents. The simulation environment is carefully constructed to foster emergent social behaviors: agents share a world, observe each other, modify their environment, and communicate directly, with resource interdependencies encouraging collaboration. Civilizational benchmarks inspired by human history are used to assess progress in role specialization, rule formation, and cultural transmission, moving beyond simple task completion metrics.

Within this framework, agents in Project Sid develop specialized roles, informal norms, and even cultural and religious transmission mechanisms without explicit programming (AL et al., 2024). Specialization arises from reinforcement and social learning, while rules and governance structures emerge through repeated social interactions and feedback. Communication patterns and task allocation strategies are not predefined but evolve as agents learn effective ways to coordinate, resulting in dynamic, adaptive social systems. These emergent behaviors include cooperation, competition, alliance formation, and conflict resolution, demonstrating how complex social dynamics can arise from simple underlying mechanisms.

Project Sid demonstrates that AI agents can autonomously achieve civiliza-

tional milestones reminiscent of human societies, including role specialization, norm creation, and cultural transmission (AL et al., 2024). However, the system faces limitations: computational constraints restrict simulation scale and duration, agent behaviors remain less complex than those of humans, and reliance on language models complicates attribution of emergent behaviors. Despite these challenges, Project Sid offers valuable insights for designing adaptable, resilient MAS and highlights open questions about long-term stability, heterogeneity, and transferability to other domains.

4.2 TeamCraft

TeamCraft is a multi-modal multi-agent benchmark in Minecraft, featuring 55,000 procedurally-generated task variants specified by visual, textual, and contextual prompts (Long et al., 2024). The benchmark is designed to evaluate collaborative capabilities by requiring agents to integrate information across modalities and coordinate on tasks ranging from simple resource gathering to complex construction. A key innovation is the use of expert demonstrations for imitation learning, with varied team sizes, resources, and environments, providing a rich dataset for training and evaluation. Evaluation protocols test generalization to novel goals, scenes, and team compositions, offering a comprehensive assessment of collaborative performance.

Agents in TeamCraft must parse visual and textual inputs to form coherent task representations, then plan and execute actions accordingly (see Figure 4) (Long et al., 2024). Coordination is achieved through explicit role assignment and task decomposition, with agents deciding how to allocate subtasks and sequence activities for efficiency. Learning is primarily imitation-based, supplemented by adaptation mechanisms for novel situations. While communication and coordination protocols are optimized for efficiency and derived from demonstrations, this structure can limit flexibility in unfamiliar contexts.

TeamCraft agents excel at efficient collaboration on familiar tasks, demonstrating strong performance when conditions match training data (Long et al., 2024). The main limitation is generalization: performance drops on novel goals, scenes, or team sizes, revealing challenges in abstracting collaboration principles beyond demonstrated examples. Potential improvements include better cross-modal reasoning, explicit communication channels, and meta-learning for adaptation. TeamCraft’s comprehensive benchmark framework advances the evaluation of collaborative MAS, but its struc-

tured approach may constrain adaptability compared to more emergent systems.

5 Comparative Analysis

5.1 Organizational Structures

Project Sid and TeamCraft represent fundamentally different approaches to organizational structure in MAS. Project Sid employs an emergent organizational model where hierarchies, roles, and relationships develop through agent interactions without predefined structures (AL et al., 2024). Agents gradually differentiate into specialized roles based on their experiences and the needs of the collective, forming flexible hierarchies that evolve with changing conditions. This approach parallels how human societies naturally organize, with leadership and role specialization emerging from repeated social interactions rather than external design.

In contrast, TeamCraft implements a more structured, task-oriented organizational approach where roles and responsibilities are explicitly defined by task requirements (Long et al., 2024). Coordination frameworks provide clear guidelines for how agents should divide labor and sequence activities, creating more predictable but less adaptable organizational patterns. This approach resembles formal organizations in human society, where roles and processes are designed to optimize efficiency for specific objectives rather than evolving organically through social dynamics.

The decision-making processes in these systems reflect their organizational philosophies. Project Sid exhibits a mix of centralized and decentralized decision-making that emerges based on task complexity and social dynamics (AL et al., 2024). Leaders naturally emerge for certain activities, while other decisions remain distributed among agents. TeamCraft, however, tends toward more systematized decision processes guided by task decomposition and role assignments derived from expert demonstrations (Long et al., 2024). These different approaches create trade offs between adaptability and efficiency, with Project Sid’s flexible decision-making allowing better adaptation to unexpected situations while TeamCraft’s structured processes optimize performance on well-defined tasks.

Each organizational approach demonstrates distinct strengths and weaknesses when adapting to changing conditions (AL et al., 2024; Chen et al., 2024; Long et al., 2024). Project Sid’s emergent structures show remarkable adaptability to environmental changes and novel social configurations, allowing the system to reorganize in

response to new challenges without external intervention. However, this adaptability comes at the cost of efficiency, as emergent organization requires time to stabilize and may not optimize for specific task performance. TeamCraft’s structured approach enables rapid mobilization for familiar tasks but struggles with novel contexts that require organizational reconfiguration, showing diminished performance when faced with unexpected challenges or team composition changes.

5.2 Communication and Coordination

Communication mechanisms differ significantly between Project Sid and TeamCraft, reflecting their contrasting organizational philosophies. Project Sid allows communication patterns to emerge naturally through agent interactions, with both direct communication (explicit message passing) and indirect communication (environmental modifications) developing as agents learn effective interaction strategies (AL et al., 2024). This emergent communication evolves to meet the needs of the collective, becoming more structured and specialized as social complexity increases.

TeamCraft, in contrast, implements more formalized communication protocols derived from expert demonstrations, with clear patterns for information sharing and coordination signaling (Long et al., 2024; Qiu et al., 2024). These protocols are optimized for task efficiency, ensuring that agents share precisely the information needed for successful collaboration without unnecessary communication overhead. While effective for familiar tasks, this approach may limit adaptability when novel situations require communication patterns not present in the training demonstrations.

Task allocation strategies similarly reflect the systems’ divergent approaches to organization. Project Sid develops allocation mechanisms through trial and error as agents learn which distributions of responsibilities lead to collective success (AL et al., 2024). This emergent allocation adapts dynamically to changing circumstances and agent capabilities, allowing the system to reorganize in response to new challenges. TeamCraft implements more explicit allocation strategies based on task structure analysis, optimizing assignments to minimize completion time and resource usage (Dong et al., 2024; Long et al., 2024). This approach yields high efficiency for well-understood tasks but may struggle with novel situations requiring creative reallocation of responsibilities.

The different communication and coordination approaches create fundamental

trade-offs in system performance (AL et al., 2024; Dong et al., 2024; Long et al., 2024; Qiu et al., 2024). Project Sid’s emergent strategies develop greater robustness to unexpected changes and novel situations, as the system can adapt its communication and coordination patterns through continued learning. However, this adaptability comes at the cost of initial efficiency, as effective patterns must evolve through experience rather than being optimized from the start. TeamCraft’s structured approaches deliver superior performance on familiar tasks through optimized protocols but show limited flexibility when facing novel challenges that require communication or coordination strategies not present in training data.

5.3 Learning and Adaptation

The learning mechanisms employed by Project Sid and TeamCraft represent distinct approaches to knowledge acquisition and behavioral adaptation. Project Sid primarily utilizes reinforcement learning combined with social learning, allowing agents to discover effective behaviors through environmental feedback while also learning from observing other agents (AL et al., 2024). This approach enables open-ended learning where new capabilities can emerge without being explicitly encoded in training data, facilitating the development of novel social structures and interaction patterns.

TeamCraft relies more heavily on imitation learning from expert demonstrations, supplemented with fine-tuning through reinforcement learning (Long et al., 2024; Zhao et al., 2024). This approach allows agents to rapidly acquire effective coordination strategies by observing successful examples, learning both task-specific knowledge and generalizable collaboration principles. The structured learning process produces more predictable performance on tasks similar to the demonstrations but may limit discovery of novel strategies that weren’t present in the training data.

Adaptation to novel situations reveals significant differences between the two systems. Project Sid demonstrates stronger adaptation to completely novel scenarios due to its emphasis on foundational learning mechanisms rather than task-specific knowledge (AL et al., 2024). Agents can apply basic principles of social organization and resource management to unfamiliar challenges, gradually developing effective responses through continued interaction and feedback. TeamCraft shows better immediate performance on variations of familiar tasks but struggles with novel challenges that require deviation from demonstrated behaviors (Long et al., 2024).

Both systems face challenges in knowledge transfer across domains and tasks, though for different reasons (AL et al., 2024; Long et al., 2024; Zhao et al., 2024). Project Sid’s emergent knowledge is often implicit in agent behaviors rather than explicitly represented, making it difficult to transfer learning directly from one context to another. TeamCraft’s more structured representations enable clearer knowledge transfer within similar task families but may create overfitting to particular task structures that hinder transfer to significantly different domains. These different learning approaches highlight the tension between flexibility and efficiency in knowledge acquisition for MAS.

5.4 Sustainability and Scalability

The long-term viability of agent collaborations differs between Project Sid and TeamCraft. Project Sid demonstrates superior sustainability through its self-organizing social structures, which continuously adapt to changing conditions and can recover from perturbations without external intervention (AL et al., 2024; Yu et al., 2024). The development of cultural transmission mechanisms allows knowledge to persist across time, while emergent governance structures help maintain social cohesion even as individual agents come and go. These features create resilient social systems capable of enduring over extended periods despite changing environmental conditions.

TeamCraft shows stronger performance in short-term task execution but may face challenges in long-term sustainability due to its reliance on predefined coordination frameworks (Long et al., 2024). Without continuous adaptation mechanisms, these frameworks may become less effective as conditions drift from their original design parameters. The system lacks the self-organizing capabilities that would allow it to reorganize in response to significant environmental changes or agent population shifts, potentially leading to performance degradation over extended periods.

Scaling to larger agent populations presents distinct challenges for each system. Project Sid demonstrates impressive scalability, successfully operating with populations ranging from 10 to over 1000 agents (AL et al., 2024). The emergent organizational structures naturally accommodate growing populations through increasing specialization and hierarchical organization, mirroring how human societies scale. This natural scalability comes at the cost of increasing organizational complexity and potential inefficiencies as coordination requirements grow more complex.

TeamCraft faces greater challenges in scaling beyond the team sizes present in its training data (Long et al., 2024; Yu et al., 2024). Performance tends to degrade when the number of agents increases significantly beyond what was demonstrated, suggesting limitations in how the coordination frameworks generalize to larger teams. This scaling challenge highlights a broader issue with task-oriented approaches: their optimization for specific contexts may create implicit assumptions about team size and composition that limit generalization to significantly different configurations.

Resource efficiency presents different trade-offs in the two systems. Project Sid initially shows lower efficiency as social structures are forming but develops increasingly efficient resource utilization as specialization and coordination mechanisms mature (AL et al., 2024). TeamCraft demonstrates higher immediate efficiency on familiar tasks but may use resources suboptimally when facing novel challenges that require adaptation (Long et al., 2024). These patterns highlight the fundamental tension between optimization for known conditions versus adaptability to changing circumstances—a key consideration for sustainable MAS.

6 Discussion

The comparison of Project Sid and TeamCraft (see Table 2) yields key insights for MAS theory and practice, particularly regarding the relationship between organizational structure and adaptive capacity. Project Sid’s emergent, self-organizing approach demonstrates that complex coordination can arise without explicit design, supporting theories of emergent complexity in MAS and highlighting the value of structural flexibility for developing adaptive collaborative systems (AL et al., 2024). In contrast, TeamCraft’s structured coordination frameworks optimize efficiency for well-defined tasks but may constrain adaptability in dynamic or unfamiliar contexts (Long et al., 2024).

These findings reveal parallels with human behavior: Project Sid’s emergent leadership and role specialization mirror organic development in human groups, while TeamCraft’s formalized processes resemble designed organizations. The results suggest that matching organizational structure to environmental uncertainty is as relevant for artificial societies as for human ones (Chen et al., 2024; Lhaksmana et al., 2018). Moreover, the accelerated and computational nature of agent societies offers a novel lens for studying social organization and value alignment in MAS.

Several design principles for effective MAS emerge from this comparison. First, combining structured coordination frameworks with emergent organizational capabilities enables systems to balance efficiency and adaptability. Second, robust communication protocols that optimize information sharing while minimizing overhead are essential for coordination across varying contexts. Third, supporting both role flexibility and specialization allows systems to develop expertise while remaining responsive to change. Fourth, mechanisms for cultural transmission facilitate knowledge preservation and system-wide adaptation over time. Finally, ethical safeguards—such as transparency, accountability, and human oversight—must be integrated from the start, especially as MAS scale and develop autonomous organizational capabilities.

The practical implications extend beyond Minecraft-based MAS. Hybrid approaches that combine efficiency with adaptability could benefit domains such as disaster response, warehouse robotics, and online community management. However, the transferability of these insights requires careful analysis of domain-specific coordination challenges and organizational requirements.

7 Conclusion and Future Work

This comparative analysis reveals fundamental differences in how MAS approach collaboration in Minecraft environments. Project Sid’s emergent, socially-oriented approach produces adaptable, self-organizing agent societies, while TeamCraft’s structured, task-oriented approach enables efficient coordination for specific challenges (AL et al., 2024; Long et al., 2024). Project Sid excels in adaptability and long-term sustainability; TeamCraft demonstrates superior efficiency on well-defined tasks. The optimal approach depends on the collaborative context: dynamic, unpredictable environments favor emergent organization, while stable, well-defined domains benefit from structured coordination.

This study is limited by its reliance on published descriptions rather than direct experimentation, differences in system objectives, and the recency of both systems. Future research should explore hybrid approaches that combine the strengths of both paradigms, conduct longitudinal studies of agent societies, investigate integration with human teams, and examine the transfer of Minecraft-based insights to real-world domains. Emphasizing the identified design principles will be crucial for developing sustainable, adaptable, and ethically aligned MAS.

References

- AL, A., Ahn, A., Becker, N., Carroll, S., Christie, N., Cortes, M., Demirci, A., Du, M., Li, F., Luo, S., Wang, P. Y., Willows, M., Yang, F., & Yang, G. R. (2024). Project Sid: Many-agent simulations toward AI civilization. <https://arxiv.org/abs/2411.00114>
- Chen, J., Jiang, Y., Lu, J., & Zhang, L. (2024). S-Agents: Self-organizing Agents in Open-ended Environments. *ICLR 2024 Workshop on Large Language Model (LLM) Agents*. <https://doi.org/10.48550/arXiv.2402.04578>
- Dong, Y., Zhu, X., Pan, Z., Zhu, L., & Yang, Y. (2024, August). VillagerAgent: A Graph-Based Multi-Agent Framework for Coordinating Complex Task Dependencies in Minecraft. In L.-W. Ku, A. Martins, & V. Srikumar (Eds.), *Findings of the association for computational linguistics: Acl 2024* (pp. 16290–16314). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2024.findings-acl.964>
- Eisenhardt, K. M. (1989). *Building theories from case study research* (Vol. 14). Academy of Management Review.
- Fan, L., Wang, G., Jiang, Y., Mandlekar, A., Yang, Y., Zhu, H., Tang, A., Huang, D.-A., Zhu, Y., & Anandkumar, A. (2022). MineDojo: Building Open-Ended Embodied Agents with Internet-Scale Knowledge. *Thirty-sixth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*. <https://doi.org/10.48550/arXiv.2206.08853>
- Gong, R., Huang, Q., Ma, X., Vo, H., Durante, Z., Noda, Y., Zheng, Z., Zhu, S.-C., & Terzopoulos, D. (2023). MindAgent: Emergent Gaming Interaction. <https://doi.org/10.48550/arXiv.2309.09971>
- Lhaksmana, K. M., Murakami, Y., & Ishida, T. (2018). Role-Based Modeling for Designing Agent Behavior in Self-Organizing Multi-Agent Systems. *International Journal of Software Engineering and Knowledge Engineering*, 28(01), 79–96. <https://doi.org/10.1142/S0218194018500043>
- Lică, M., Shirekar, O., Colle, B., & Raman, C. (2024). MindForge: Empowering Embodied Agents with Theory of Mind for Lifelong Cultural Learning. <https://doi.org/10.48550/arXiv.2411.12977>

- Long, Q., Li, Z., Gong, R., Wu, Y. N., Terzopoulos, D., & Gao, X. (2024). TeamCraft: A Benchmark for Multi-Modal Multi-Agent Systems in Minecraft. <https://arxiv.org/abs/2412.05255>
- Lund, B., Wang, X., Wang, Y., et al. (2023). ChatGPT and artificial intelligence in higher education: A tsunami of questions and a sea of opportunities. *International Journal of Educational Technology in Higher Education*, 20(1), 1–24. <https://doi.org/10.1186/s41239-023-00444-9>
- Panait, L., & Luke, S. (2005). Cooperative Multi-Agent Learning: The State of the Art. *Autonomous Agents and Multi-Agent Systems*, 11(3), 387–434. <https://doi.org/10.1007/s10458-005-2631-2>
- Qiu, X., Wang, H., Tan, X., Qu, C., Xiong, Y., Cheng, Y., Xu, Y., Chu, W., & Qi, Y. (2024). Towards Collaborative Intelligence: Propagating Intentions and Reasoning for Multi-Agent Coordination with Large Language Models. *CoRR*, *abs/2407.12532*. <https://doi.org/10.48550/arXiv.2407.12532>
- Russell, S., & Norvig, P. (2009). *Artificial Intelligence: A Modern Approach* (3rd). Prentice Hall Press. <https://doi.org/10.5555/1671238>
- Shu, R., Das, N., Yuan, M., Sunkara, M., & Zhang, Y. (2024). Towards Effective GenAI Multi-Agent Collaboration: Design and Evaluation for Enterprise Applications. <https://arxiv.org/abs/2412.05449>
- Smyth, E., & Suglia, A. (2025). VoyagerVision: Investigating the Role of Multi-modal Information for Open-ended Learning Systems. <https://doi.org/10.48550/arXiv.2507.00079>
- van Leeuwen, C. J. (2021). *Self-Organizing Multi-Agent Systems* [Doctoral thesis]. Delft University of Technology. <https://doi.org/10.4233/uuid:face1276-5025-4853-996a-30d6771cf24d>
- Wang, Z., Cai, S., Liu, A., Jin, Y., Hou, J., Zhang, B., Lin, H., He, Z., Zheng, Z., Yang, Y., Ma, X., & Liang, Y. (2025). JARVIS-1: Open-World Multi-Task Agents With Memory-Augmented Multimodal Language Models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(3), 1894–1907. <https://doi.org/10.1109/TPAMI.2024.3511593>
- Wooldridge, M. (2009). *An Introduction to MultiAgent Systems* (2nd). Wiley Publishing. <https://doi.org/10.5555/1695886>

- Yin, R. K. (2018). *Case Study Research and Applications: Design and Methods* (6th). SAGE Publications.
- Yu, X., Fu, J., Deng, R., & Han, W. (2024). MineLand: Simulating Large-Scale Multi-Agent Interactions with Limited Multimodal Senses and Physical Needs. <https://arxiv.org/abs/2403.19267>
- Zhao, Z., Chai, W., Wang, X., Ma, K., Chen, K., Guo, D., Ye, T., Zhang, Y., Wang, H., & Wang, G. (2024). STEVE Series: Step-by-Step Construction of Agent Systems in Minecraft. *CoRR*, *abs/2406.11247*. <https://doi.org/10.48550/arXiv.2406.11247>

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

A Figures and Tables




Figure 1: An example Minecraft technology dependency tree for the mining of gold, diamond, and emeralds (AL et al., 2024).




Figure 2: Illustration of trading activities among agents in a Minecraft multi-agent system (AL et al., 2024).



Figure 3: Visual abstract of Project Sid, illustrating the progression from agent architecture to emergent agent civilization (AL et al., 2024).

System Prompt	Three Orthographic Views	
	Language Instruction	Three bots need to build a building on the platform. Write actions for <i>bot1</i> , <i>bot2</i> , <i>bot3</i> based on given observation .
Observation	First Person View	
	Inventory Information	bot1 has 5 bricks, 3 iron_ore, bot2 has 2 sea_lantern, bot3 has 1 brick...
Action		<code>placeItem(bot1, 'bricks', (-1,0,-1))</code> <code>placeItem(bot2, 'oak_planks', (0,0,0))</code> <code>placeItem(bot3, 'iron_ore', (0,0,-1))</code>

Figure 4: Example of a multi-modal prompt for building in Minecraft (Long et al., 2024).

Table 1: Comparison of candidate case studies with features and selection rationale.

Case Study	Year	Collaboration Paradigm	Documentation Quality	Distinctiveness	Selection	Rationale
Project Sid	2024	Emergent/Social	High	High	Included	Most robust, clear emergent paradigm
TeamCraft	2024	Structured/Task	High	High	Included	Most robust, clear structured paradigm
S-Agents	2024	Emergent/Social (with task features)	High	Moderate	Excluded	Overlaps with Project Sid, not a distinct paradigm
MineDojo	2022	Task-oriented	High	Low	Excluded	Single-agent focus, not multi-agent collaboration
JARVIS-1	2025	Task-oriented	High	Low	Excluded	Single-agent focus, not multi-agent collaboration
VillagerAgent	2024	Structured/Task (Graph-based)	Moderate	Moderate	Excluded	Overlaps with TeamCraft, not a distinct paradigm
MineLand	2024	Emergent/Social (Ecological)	Moderate	Moderate	Excluded	Focuses on ecology/social simulation, not collaboration paradigm
STEVE Series	2024	Single-agent/ Task-oriented	High	Low	Excluded	Not multi-agent collaboration

Continued on next page

Continued from previous page

Case Study	Year	Collaboration Paradigm	Documentation Quality	Distinctiveness	Selection	Rationale
VoyagerVision	2025	Single-agent/ Task-oriented	High	Low	Excluded	Not multi-agent collaboration
MindForge	2024	Platform (multi-agent support)	High	Moderate	Excluded	Infrastructure only, not a novel paradigm
MindAgent	2023	Single-agent/ Task-oriented	High	Low	Excluded	Not multi-agent collaboration

Table 2: Comparison of Project Sid and TeamCraft across key dimensions.

Dimension	Project Sid	TeamCraft	Comparison
Organizational Structure	Emergent, self-organizing; roles and hierarchies develop dynamically; flexible and adaptive to change.	Structured, task-oriented; roles and responsibilities predefined by task requirements; optimized for efficiency.	Project Sid: Highly adaptive, mirrors organic societies. TeamCraft: Predictable, efficient for set tasks.
Communication and Coordination	Communication patterns and coordination mechanisms emerge through agent interaction; both direct and indirect (environmental) communication.	Formalized protocols derived from demonstrations; explicit, efficient information sharing for task completion.	Project Sid: Evolves communication, robust to novelty. TeamCraft: Optimized protocols, excels in routine.
Learning and Adaptation	Reinforcement and social learning; open-ended, agents learn from environment and each other; supports novel behavior emergence.	Imitation learning from expert demonstrations, with some reinforcement learning; strong on familiar tasks, limited on novel ones.	Project Sid: Open-ended, adapts to new scenarios. TeamCraft: Fast on known tasks, less generalization.

Continued on next page

Continued from previous page

Dimension	Project Sid	TeamCraft	Comparison
Sustainability and Scalability	Self-organizing, resilient social structures; adapts to agent turnover and environmental change; scales to 1000+ agents.	Predefined frameworks; strong short-term performance, but less sustainable as conditions change; scaling beyond training data is challenging.	Project Sid: Long-term, scalable, resilient. TeamCraft: High efficiency, but scaling and sustainability limited.

B AI Conversation Logs

```
1 <?xml version="1.0" encoding="UTF-8"?>
2 <conversationHistory>
3   <tool>Perplexity Research</tool>
4   <model>o4-mini</model>
5   <exchange id="1">
6     <user>
7       I'm a Master's student in Information Systems. You'll act
      as my AI research assistant. Together we will draft a 15-page
      scientific paper titled "When Blocks Build Societies: Multi-
      Agent Systems in Minecraft." We will compare Project Sid and
      TeamCraft, analyze their approaches to sustainable multi-agent
      collaboration, and document every step of the process for
      transparency.
8       Let's begin by defining the scope and RQ precisely.
9     </user>
10    <agent>
11      Understood. Proposed research question:
12      "In Minecraft-based multi-agent environments, how do
      emergent social structures (Project Sid) contrast with
      explicit coordination protocols (TeamCraft) in terms of
      sustainability, scalability, and adaptability of collaboration
      "?
13      Scope: compare architectural designs, communication
      mechanisms, learning paradigms, evaluation metrics, and long-
      term viability.
14      Next, shall I draft a detailed paper outline with section
      lengths?
15    </agent>
16  </exchange>
17
18  <exchange id="2">
19    <user>
20      Yes. Provide a 15-page outline, indicating sections,
      subsections, and approximate page allocation.
21    </user>
22    <agent>
23      Here's a draft outline:
```

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

```

24     1. Introduction (1p)
25     2. Background & Related Work (2.5p)
26         2.1 Multi-Agent Systems (1p)
27         2.2 Minecraft as a Testbed (0.75p)
28         2.3 Collaboration Paradigms (0.75p)
29     3. Methodology (1.5p)
30         3.1 Case Study Selection (0.5p)
31         3.2 Analytical Framework (0.5p)
32         3.3 Use of Generative AI (0.5p)
33     4. Case Study 1: Project Sid (2.5p)
34         4.1 Architecture & Setup (1p)
35         4.2 Emergent Social Systems (1p)
36         4.3 Findings & Limitations (0.5p)
37     5. Case Study 2: TeamCraft (2.5p)
38         5.1 Benchmark Design (1p)
39         5.2 Coordination Framework (1p)
40         5.3 Findings & Limitations (0.5p)
41     6. Comparative Analysis (2p)
42         6.1 Organizational Structures (0.5p)
43         6.2 Communication & Coordination (0.5p)
44         6.3 Learning & Adaptation (0.5p)
45         6.4 Sustainability & Scalability (0.5p)
46     7. Discussion (1p)
47     8. Conclusion & Future Work (0.5p)
48     9. References & Appendix (1p)
49     Shall we proceed to flesh out the Background section?
50     </agent>
51 </exchange>
52
53 <exchange id="3">
54     <user>
55         Yes. Expand Background & Related Work. Include literature
56         on MineDojo, JARVIS-1, and theories of emergent organization
57         vs. structured coordination.
58     </user>
59     <agent>
60         Draft Background content:
61         • Multi-Agent Systems: evolution from distributed AI to
62         deep RL \citep{stone2000multiagent, wooldridge2009introduction

```

```

    }.
60     • Minecraft as Testbed: MineDojo framework \citep{
fan2022minedojo}, JARVIS-1 multimodal LLM \citep{
zhou2023jarvis}.
61     • Collaboration Paradigms: emergent coordination \citep{
leibo2019autocurricula}, task-oriented protocols \citep{
yang2020survey}.
62     I'll integrate citations in APA style and elaborate on each
subtopic. Ready for Methodology?
63     </agent>
64 </exchange>
65
66 <exchange id="4">
67     <user>
68         Go ahead with Methodology. Make sure to describe our use of
generative AI for literature search, drafting, and
verification.
69     </user>
70     <agent>
71         Draft Methodology:
72         - Case Study Selection: rationale for choosing Sid and
TeamCraft (recency, contrasting approaches).
73         - Analytical Framework: four dimensions (structure,
communication, learning, sustainability) with metrics from MAS
and org theory.
74         - Generative AI Usage: tools (GPT-4, Perplexity), prompts
used for literature identification, synthesis, outline
generation, drafting.
75         - Ethical & Transparency Measures: human oversight,
citation verification, full conversation logs in appendix.
76         Shall I proceed to flesh out the Project Sid section?
77     </agent>
78 </exchange>
79
80 <exchange id="5">
81     <user>
82         Yes. For Case Study 1, draft the "Architecture & "Setup and
"Emergent Social "Systems subsections with citations.
83     </user>

```

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

```

84     <agent>
85         Case Study 1 draft:
86         4.1 Architecture & Setup:
87             - PIANO architecture enabling multimodal perception and
real-time orchestration \citep{wang2024sid}.
88             - Simulation parameters: -101000 agents, resource rules,
communication channels.
89         4.2 Emergent Social Systems:
90             - Role specialization via reinforcement & social learning.
91             - Rule formation: norms to formal governance.
92             - Cultural transmission: shared beliefs & rituals.
93         Cite examples and quantitative benchmarks. Ready to draft
Findings & Limitations?
94     </agent>
95 </exchange>
96
97 <exchange id="6">
98     <user>
99         Yes. Then move on to Case Study 2: "Benchmark "Design and
"Coordination Framework".
100    </user>
101    <agent>
102        Case Study 2 draft:
103        5.1 Benchmark Design:
104            - 55,000 multi-modal task variants (visual + textual
prompts) \citep{li2024teamcraft}.
105            - Procedural generation of expert demonstrations for
imitation learning.
106        5.2 Coordination Framework:
107            - Role assignment, task decomposition, communication
protocols.
108            - Evaluation on generalization: novel goals, scenes, team
compositions.
109        Include performance metrics and observed failure modes.
Shall I write Findings & Limitations next?
110    </agent>
111 </exchange>
112
113 <exchange id="7">

```

```

114     <user>
115         Yes. Then draft Comparative Analysis: structure it by the
116         four dimensions and include brief comparative comments under
117         each.
118     </user>
119     <agent>
120         Comparative Analysis draft:
121         6.1 Organizational Structures: Sid's emergent hierarchies
122         vs. TeamCraft's explicit roles.
123         6.2 Communication & Coordination: emergent messaging/
124         stigmergy vs. fixed protocols.
125         6.3 Learning & Adaptation: RL + social learning vs.
126         imitation + fine-tuning.
127         6.4 Sustainability & Scalability: Sid's resilience &
128         cultural memory vs. TeamCraft's short-term efficiency.
129         For each, include trade-offs and cite case study results.
130         Proceed to Discussion?
131     </agent>
132 </exchange>
133
134 <exchange id="8">
135     <user>
136         Yes .Write the Discussion and Conclusion & Future Work
137         sections, highlighting theoretical and practical implications.
138     </user>
139     <agent>
140         Discussion draft:
141         - Theoretical: emergent complexity theory, parallels to
142         human organizational behavior, philosophical questions on AI
143         autonomy.
144         - Practical: hybrid system design, applications in robotics
145         /logistics, ethical considerations (transparency, oversight).
146         Conclusion & Future Work draft:
147         - Summary of Sid vs. TeamCraft strengths/limitations.
148         - Proposals: hybrid architectures, longitudinal studies,
149         human-AI integration, domain transfer.
150         Ready for LaTeX skeleton and bibliography?
151     </agent>
152 </exchange>

```

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

```

141
142 <exchange id="9">
143   <user>
144     Yes. Generate a LaTeX skeleton ('main.tex') with sections,
        cite commands ('\citep{}'), and placeholders for tables/
        figures.
145   </user>
146   <agent>
147     LaTeX skeleton prepared, using 'natbib' and 'apa' style,
        section headings matching outline. Places '\bibliography{
        literature}'.
148     I will include comment markers ('% TODO:') for figures and
        tables.
149     Shall I draft the 'literature.bib' file next?
150   </agent>
151 </exchange>
152
153 <exchange id="10">
154   <user>
155     Please create 'literature.bib' with all cited sources in
        APA format.
156   </user>
157   <agent>
158     Generated 'literature.bib' entries for Wooldridge (2009),
        Stone \& Veloso (2000), Fan et al. (2022), Zhou et al. (2023),
        Wang et al. (2024), Li et al. (2024), Yang et al. (2020),
        Leibo et al. (2019), Lund et al. (2023).
159     All entries include DOI/URL where available.
160     Next: ensure in-text citations follow latest APA guidelines
        ?
161   </agent>
162 </exchange>
163
164 <exchange id="11">
165   <user>
166     Yes. Review in-text citations and reference formatting to
        match APA 7th edition.
167   </user>
168   <agent>

```

169 Updated in-text citations to author-year parenthetical
style, ensured reference list entries include proper title
capitalization, journal italics, DOIs.

170 Verified consistency with APA 7. All good.

171 Finally, export this entire conversation as XML for the
appendix.

172 </agent>

173 </exchange>

174

175 <exchange id="12">

176 <user>

177 Perfect. Provide the conversation history as structured XML
.

178 </user>


179 <agent>

180 Provided above. This XML log captures each iteration
transparently, fulfilling the paper's methodology requirement
for AI usage disclosure.

181 </agent>

182 </exchange>

183 </conversationHistory>

University
of Cologne

in seminar papers, term papers, bachelor theses, and master theses

Subtitle

		Usage	
Tools		Yes	No
a.	ChatGPT	☐	☐
b.	DeepL Write	☐	☐
c.	Microsoft 365 Copilot	☐	☐
d.	Neuroflash	☐	☐
e.	Grammarly	☐	☐
f.	DeepL Translator	☐	☐
g.	Google Translate	☐	☐
h.	Perplexity	☐	☐
i.	Elicit	☐	☐
j.	Explainpaper	☐	☐
k.	ResearchRabbit	☐	☐
l.	GitHub Copilot	☐	☐
m.	Other tool	☐	☐
n.	Other tool	☐	☐
o.	Other tool	☐	☐

Statements on responsible tool use	Confirmation: Yes
I/we are informed about the capabilities and limitations of the tools I/we used.	☐
I/we verified that the results given by the tools are accurate.	☐
I/we acknowledge that the responsibility for the paper/thesis lies with the author(s), not the tools or anybody else.	☐

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

3. Detailed activities for which the tools were used

Indicate which tool you used and to what extent for which activity.

Activity	Description	Tools used (if any) ¹	Description of the tool usage (type of usage, affected sections of the paper/thesis, etc.)
Ideation & conceptualization	- Generating ideas, re-search goals, aims, and questions - Identifying and defining relevant concepts		
Literature search and analysis	- Searching for relevant literature - Reviewing potentially relevant literature - Summaries of relevant literature		
Methodology	- Searching for an appropriate methodology - Designing and tailoring the methodology to the		
Coding	- Creating and documenting code, algorithms, software - Testing and debugging of existing code, algorithms, software - Understanding existing code, algorithms, software		
Data collection and analysis	- Collection of primary or secondary data - Qualitative data analysis (including summarizing and coding) - Quantitative data analysis (including statistics) - Mathematical, computational, or other formal techniques for modeling, simulation, and analytics		
Interpretation and Validation	- Interpretation of results - Derivation of implications for research and practice - Verification of the overall replication/ reproducibility of results and other re-search outputs		
Structuring and planning the text	- Outlining the paper/thesis - Outlining sections of the paper/thesis (e.g., bullet point lists per section)		

¹ Provide letters from the tool list in section 1.

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

Generating the text	– Generating text on various topics in different sections of the paper (including title and abstract)		
Translating text	– Translating text written by the authors – Translating text written by others		
Reviewing & editing the text	– Critical review, feedback or revision on content, organization, or grammar of the paper – Proofreading – Rephrasing or paraphrasing text – Shortening / extending text		
Presentation	– Structuring a presentation on the paper – Filling a presentation with content on the paper		
Citation Management	– Creation of reference list – Formatting of references		
Further activities			

Is there anything else to declare regarding the use of AI tools that would highlight or limit your independent, definable performance in crafting this paper/thesis?

DRAFT-9feb662 @ 2025-07-29 19:40:25 +0200

4. Signature(s) of author(s)

	Author last name, first name(s)	Matriculation number	Date	Signature
1				

Background on this declaration: Assessments at the university must allow students' performance to be determined. For this purpose, the student's own performance, which is contained in a paper/thesis, must be distinguished from the performance of others. It is also necessary to understand the tools used. As long as generative AI is relatively new in higher education, its use must be declared in detail. This serves the evaluation of the performance and assures students that their work will not be reassessed afterwards against the background of other norms.

Sources: Gimpel, H., Dilger, P. (2023). Declaration of Generative AI Tool Usage in seminar papers, term papers, bachelor theses, and master theses. University of Hohenheim, May 3, 2023.