

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

When Blocks Build Societies
Multi-Agent Systems in Minecraft

Jan-Philipp Kiel
University of Cologne

Advanced Seminar
Information Systems and Digital Technology
Prof. Dr. Stefan Seidel

July 28, 2025

Table of Contents

1	Introduction.....	1
2	Background and Related Work.....	2
2.1	Multi-Agent Systems in Virtual Environments	2
2.2	Minecraft as a Research Platform.....	3
2.3	Approaches to Agent Collaboration.....	4
3	Methodology	5
3.1	Case Study Selection.....	6
3.2	Literature Search and Screening	7
3.3	Analytical Framework.....	7
3.4	Use of Generative AI	8
4	Case Study 1: Project Sid	9
4.1	Self-Organizing Social Systems.....	10
4.2	Key Findings and Limitations	11
5	Case Study 2: TeamCraft.....	12
5.1	Task-Oriented Coordination Frameworks.....	13
5.2	Key Findings and Limitations	14
6	Comparative Analysis	15
6.1	Organizational Structures	15
6.2	Communication and Coordination.....	16
6.3	Learning and Adaptation.....	17
6.4	Sustainability and Scalability	19
7	Discussion.....	20
7.1	Theoretical Implications	20
7.2	Practical Applications	21
8	Conclusion and Future Work	22
	References	24
A	Comparison of Candidate Case Studies.....	27
B	AI Conversation History	27
	Declaration of Generative AI Tool Usage.....	35

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

1 Introduction

Multi-agent systems (MAS) are a key frontier in artificial intelligence (AI) research, offering insights into how autonomous entities can collaborate, compete, and coexist within shared environments. These systems enable the study of complex emergent behaviors and collective intelligence that extend beyond the capabilities of single agents (Shu et al., 2024; van Leeuwen, 2021). As AI continues to advance, understanding the principles that govern effective multi-agent collaboration becomes increasingly vital for applications ranging from robotic teams to virtual assistants.

Despite advances in MAS, there is limited understanding of how different collaborative frameworks—task-oriented versus socially-oriented—impact sustainable collaboration in complex virtual environments such as Minecraft. This paper addresses the need to compare these approaches to inform future design of collaborative MAS. The scope of this research is limited to Minecraft-based MAS. While potential parallels to other domains are discussed, all findings and design principles are grounded in the Minecraft context.

Minecraft, with its open-world sandbox, procedurally generated environments, and flexible crafting, make it an ideal platform for multi-agent research (Fan et al., 2022). Its block-based world allows for precise control while still presenting meaningful challenges in navigation, resource gathering, construction, and social interaction (Zhao et al., 2024).

Most MAS research in Minecraft has focused on specific task completion, such as navigation or building predefined structures. However, a significant gap exists between these narrowly defined task-oriented systems and more complex, socially-oriented agent societies (Qiu et al., 2024; Shu et al., 2024). Bridging this gap requires new approaches to agent architecture, communication protocols, and coordination mechanisms that can support sustainable collaboration over extended periods.

Two recent research initiatives have made significant progress in addressing different aspects of this challenge: Project Sid and TeamCraft. Project Sid, introduced in 2024, explores large-scale simulations of 10-1000+ AI agents, focusing on how these agents develop specialized roles, adhere to collective rules, and engage in cultural transmission within a civilization-inspired framework (Alterra et al., 2024). In contrast, TeamCraft provides a multi-modal multi-agent benchmark featuring 55,000

task variants designed to test various aspects of agent collaboration through structured prompts and demonstrations (Long et al., 2024).

This paper therefore addresses the research question: *How do different collaborative frameworks in Minecraft address the challenge of sustainable multi-agent collaboration within virtual environments?* By comparing these approaches, we aim to identify design principles that can inform the development of more robust and adaptable collaborative AI systems within virtual environments.

The paper is structured as follows: First, we review background literature on MAS, Minecraft as a research platform, and approaches to agent collaboration. Next, we detail our methodology, including case study selection criteria and analytical framework. We then analyze Project Sid and TeamCraft individually before conducting a comparative analysis across organizational structures, communication patterns, learning mechanisms, and sustainability. Finally, we discuss theoretical implications, explore speculative practical applications in other domains, and conclude with directions for future research.

2 Background and Related Work

To ensure conceptual clarity, this paper adopts the following definitions: An *agent* is an entity perceiving and acting on its environment (Russell & Norvig, 2009); *artificial intelligence (AI)* is the study of agents maximizing cumulative reward (Russell & Norvig, 2009); a *multi-agent system (MAS)* consists of multiple interacting agents whose collective behavior emerges from their interactions (Wooldridge, 2009); and *collaboration* refers to agents working together to achieve shared goals, requiring coordination and communication (Panait & Luke, 2005).

2.1 Multi-Agent Systems in Virtual Environments

Multi-agent systems (MAS) represent a paradigm in artificial intelligence research where multiple autonomous agents interact within a shared environment to achieve individual or collective goals. The field has evolved significantly from early distributed AI research in the 1980s to contemporary approaches leveraging deep learning and reinforcement learning (Shu et al., 2024; van Leeuwen, 2021). Modern MAS research encompasses diverse architectures ranging from fully cooperative to competitive and mixed-motive scenarios, with applications spanning robotics, logistics,

gaming, and social simulation.

Virtual environments offer particularly valuable platforms for multi-agent research due to their controlled conditions, cost-effectiveness, and ability to simulate complex scenarios that would be impractical or impossible in physical settings (Lhaksmana et al., 2018; Shu et al., 2024). These environments provide a “consequence-free” space to test agent behaviors that might be risky or expensive in real-world settings, while still capturing essential dynamics of multi-agent interaction. Additionally, virtual environments allow researchers to precisely control experimental conditions and collect comprehensive data about agent behaviors and interactions.

Multi-agent architectures can be broadly categorized as centralized or decentralized, with various hybrid approaches emerging to leverage the advantages of both (van Leeuwen, 2021). Centralized approaches maintain a global view of the environment and coordinate agent actions through a central controller, offering optimal decision-making at the cost of scalability and robustness. Decentralized approaches distribute decision-making among individual agents, providing resilience and scalability at the expense of potentially suboptimal coordination. Communication models within these architectures range from direct message passing to stigmergic coordination through environmental modifications.

Recent advances in multi-agent collaboration have been driven by developments in large language models (LLMs) and deep reinforcement learning (Qiu et al., 2024; Shu et al., 2024). LLMs have enhanced agents’ ability to understand complex instructions and generate contextually appropriate responses, while reinforcement learning techniques have improved adaptation to dynamic environments and coordination among heterogeneous agents (Chen et al., 2024). These advances have expanded the scope of multi-agent research to include emergent communication protocols, collective intelligence, and social learning.

2.2 Minecraft as a Research Platform

Minecraft offers a uniquely suitable platform for AI research due to its procedurally generated worlds, comprehensive crafting system, and physics simulation that strikes a balance between realism and abstraction (Fan et al., 2022). The game’s block-based structure provides a discretized state space that is more computationally tractable than continuous environments while still offering sufficient complexity to

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

model interesting real-world problems. Furthermore, its open-world nature allows for a wide range of tasks, from simple navigation to complex collaborative construction projects.

Previous AI research in Minecraft has explored diverse challenges, including single-agent task completion, navigation in procedurally generated landscapes, and creative building (Wang et al., 2025; Zhao et al., 2024). The MineDojo framework, for instance, provides thousands of diverse open-ended tasks and an internet-scale knowledge base with Minecraft videos, tutorials, and forum discussions to facilitate agent learning (Fan et al., 2022). Similarly, JARVIS-1 demonstrates how multimodal language models can perceive visual observations and textual instructions to generate sophisticated plans and perform embodied control in Minecraft (Wang et al., 2025).

The Minecraft environment presents unique challenges for AI research, including a large state space, partial observability, and the need for long-horizon planning (Fan et al., 2022; Wang et al., 2025). Agents must reason about objects that may be out of view, plan sequences of actions that may take hundreds or thousands of steps to complete, and adapt to unexpected environmental changes. These challenges make Minecraft particularly valuable for testing the limits of current AI approaches and driving innovations in planning, perception, and learning algorithms.

Beyond technical challenges, Minecraft enables social interaction and emergent behaviors that are difficult to study in more constrained environments (Altera et al., 2024; Yu et al., 2024). The MineLand simulator, for example, introduces features like limited multimodal senses and physical needs to create more ecologically valid social interactions among large groups of agents (Yu et al., 2024). This capability for modeling complex social dynamics makes Minecraft an ideal test bed for studying how artificial societies might form, evolve, and sustain themselves over time.

2.3 Approaches to Agent Collaboration

Agent collaboration approaches typically fall into two broad categories: task-oriented collaboration, which focuses on completing specific objectives, and social-oriented collaboration, which emphasizes emergent social structures and norms (Qiu et al., 2024; Shu et al., 2024). Task-oriented approaches typically define explicit goals, roles, and protocols for interaction, prioritizing efficiency and measurable outcomes. Social-oriented approaches, in contrast, allow organizational structures and behavioral

norms to emerge through repeated interactions, often leading to more adaptable but less predictable collaborative behaviors.

Coordination mechanisms in MAS include role assignment, task allocation, and conflict resolution strategies (Dong et al., 2024; Shu et al., 2024). Role-based approaches assign specific responsibilities to agents based on their capabilities or position within an organizational structure. Task allocation methods distribute work among agents to optimize resource utilization and minimize conflicts. These mechanisms may be static (predetermined) or dynamic (adapting to changing conditions), with varying degrees of centralization in decision-making authority.

Communication strategies play a crucial role in effective collaboration, ranging from direct communication through message passing to indirect communication through environmental modifications (Qiu et al., 2024; Shu et al., 2024). Direct communication allows for precise information sharing but requires protocols for when, what, and how to communicate. Indirect communication (stigmergy) relies on agents observing and responding to changes in the shared environment, often leading to emergent coordination patterns. The choice of communication strategy significantly impacts a system's scalability, robustness, and adaptability to novel situations.

Evaluation metrics for successful collaboration include task completion rates, resource efficiency, adaptability to changing conditions, and robustness to agent failures (Chen et al., 2024; Qiu et al., 2024; Shu et al., 2024). Task-oriented systems typically prioritize metrics related to goal achievement and efficiency, such as time-to-completion and resource consumption. Social-oriented systems may emphasize metrics related to sustainability and resilience, such as the stability of social structures and adaptability to unexpected challenges. These different evaluation priorities reflect fundamental differences in how collaboration is conceptualized across approaches.

3 Methodology

To ensure transparency and replicability, the methodology follows a clear sequence of steps. First, candidate systems were identified and screened according to explicit inclusion and exclusion criteria. Next, the literature search and screening process was conducted to operationalize these criteria and select the most relevant systems. Following selection, a structured analytical framework was applied to compare the chosen case studies across key dimensions. Finally, the use of generative AI tools

in the research process is described, including ethical considerations and verification procedures. This sequence ensures that each methodological decision is traceable and that the comparative analysis is grounded in a systematic and reproducible approach.

The comparative case study methodology is particularly well-suited to addressing the research question, as it enables in-depth analysis of contrasting collaborative frameworks within their real-world context, following recommendations by Yin (2018) and Eisenhardt (1989) for theory-building through case comparison (Eisenhardt, 1989; Yin, 2018). The use of a structured analytical framework ensures systematic evaluation across relevant dimensions, enhancing the rigor and transparency of the analysis (Miles & Huberman, 1994). Incorporating generative AI tools supports efficient literature synthesis and data processing, aligning with emerging best practices for AI-assisted research while maintaining human oversight and accountability (Lund et al., 2023). Together, these methodological choices provide a robust foundation for exploring how different collaborative paradigms address sustainable multi-agent collaboration in Minecraft.

Despite these strengths, the methodology is subject to certain limitations. Potential selection bias may arise from the focus on recent, well-documented case studies, and the literature search may not capture all relevant systems due to database and keyword constraints. The use of generative AI tools, while increasing efficiency, introduces risks of incomplete synthesis or propagation of inaccuracies. To mitigate these issues, multiple databases were searched, inclusion and exclusion criteria were explicitly defined, and all AI-generated content was systematically reviewed and verified against original sources. These steps help ensure the validity and transparency of the comparative analysis.

3.1 Case Study Selection

Project Sid and TeamCraft were selected as the most methodologically robust and conceptually contrasting case studies of multi-agent collaboration in Minecraft. Selection followed criteria outlined by Yin (2018) and Eisenhardt (1989) and was based on recency, documentation quality, and clear contrast in collaborative paradigms. Other modern case studies were excluded for reasons such as overlap in paradigm, lack of focus on multi-agent collaboration, or insufficient distinctiveness. A detailed overview of all candidate case studies and selection rationale is provided in Table 1.

3.2 Literature Search and Screening

The literature search for candidate case studies was conducted between June and July 2025 using major academic databases including Google Scholar, Semantic Scholar, and arXiv. Search keywords included combinations of “multi-agent systems”, “Minecraft”, “collaboration”, “benchmark”, and “emergent behavior”. Initial results were screened for relevance based on title and abstract. Full texts were reviewed for systems that focused on multi-agent collaboration in Minecraft environments. Generative AI tools were used to assist in the screening and review of identified papers, enabling efficient processing and synthesis of the literature. Additionally, conversational learning with AI was used to quickly understand key concepts and frameworks; the conversation histories are referenced in the appendix (see Appendix B).

3.3 Analytical Framework

The comparative analysis employs a structured framework examining four key dimensions: (1) Organizational structure, (2) Communication and coordination, (3) Learning and adaptation, and (4) Sustainability and scalability. Each dimension is analyzed through the lens of established MAS literature and organizational theory, providing a systematic basis for comparison.

Organizational structure is relevant because the way agents are organized—whether through emergent hierarchies or predefined roles—directly influences the system’s ability to support sustainable collaboration and adapt to changing environments. Communication and coordination are central to effective multi-agent collaboration, as the mechanisms for information exchange and joint action determine how well agents can work together to achieve shared goals. Learning and adaptation are crucial for long-term sustainability, enabling agents to improve their behaviors and strategies in response to new challenges and evolving contexts. Sustainability and scalability address the research question by evaluating whether collaborative frameworks can maintain performance and resilience as the number of agents and complexity of tasks increase, which is essential for robust multi-agent societies in virtual environments like Minecraft.

Organizational structure is analyzed through the lens of hierarchy (the levels of authority and decision-making in the system), role specialization (how agents develop specialized functions), and adaptability (how the organization responds to changing

conditions) (Chen et al., 2024; Lhaksmana et al., 2018). Project Sid’s emergent hierarchies are contrasted with TeamCraft’s task-oriented organization to understand how different structural approaches affect collaborative outcomes. This analysis draws on organizational theory concepts such as centralization, formalization, and departmentalization to characterize the systems’ organizational features.

Agent autonomy and communication patterns are evaluated by examining communication frequency (how often agents exchange information), content (what information is shared), and effectiveness (how communication impacts task performance) (Qiu et al., 2024; Shu et al., 2024). The analysis considers both direct communication through message passing and indirect communication through environmental modifications. Particular attention is paid to how communication patterns support or hinder coordination in different contexts, such as novel tasks or unexpected environmental changes.

Task allocation and sustainability are measured through metrics including efficiency (resource usage relative to task completion), adaptability (performance on novel tasks or with novel team compositions), and long-term viability (stability of performance over extended periods) (Chen et al., 2024; Shu et al., 2024). These metrics provide insight into how well each system balances short-term performance goals with long-term sustainability concerns, a critical consideration for enduring multi-agent collaborations.

3.4 Use of Generative AI

This research employed generative AI tools throughout the research process, including literature review, analysis, and manuscript preparation. Specifically, LLMs were used to assist in identifying relevant literature, synthesizing findings across sources, and generating initial drafts of paper sections. This approach allowed for efficient processing of the extensive literature on MAS while maintaining focus on the specific research question.

The primary AI tool used was a LLM with research capabilities, which provided assistance in summarizing academic papers, identifying connections between concepts, and generating coherent prose based on specified content requirements. The model was provided with search results from academic databases and asked to summarize key findings, compare approaches, and identify research gaps. Human oversight

was maintained throughout this process, with all AI-generated content reviewed, revised, and expanded based on direct examination of the original sources.

Ethical considerations in the use of generative AI included ensuring proper attribution of ideas to original sources, verifying factual claims against primary literature, and maintaining transparency about the role of AI in the research process. All content generated by AI was treated as a first draft requiring human verification and refinement rather than as finished work. This approach aligns with emerging best practices for AI-assisted academic writing, which emphasize human oversight and accountability.

The use of generative AI introduces certain limitations and potential biases, including the risk of propagating inaccuracies present in the training data, generating plausible-sounding but incorrect information, and prioritizing certain types of sources over others. To mitigate these risks, a systematic verification process was implemented, comparing AI-generated content against original sources and seeking additional sources where information appeared incomplete or potentially biased. The complete record of interactions with the AI system is included in Appendix B to ensure full transparency about the research process. All sources were properly cited, and all AI-generated content was reviewed and verified by the author to ensure accuracy and compliance with ethical standards for academic research and responsible AI use.

4 Case Study 1: Project Sid

Project Sid represents a pioneering effort to study large-scale multi-agent interactions that mirror the full spectrum of civilizational processes. This research initiative simulates societies of 10 to over 1000 AI agents within a Minecraft environment, focusing on how these agents develop complex social structures and behaviors through extended interaction (Altera et al., 2024). The project introduces the PIANO (Parallel Information Aggregation via Neural Orchestration) architecture, which enables agents to interact with humans and other agents in real-time while maintaining coherence across multiple output streams.

The PIANO architecture forms the technical foundation of Project Sid, providing a framework for agent perception, reasoning, and communication (Altera et al., 2024). This architecture allows agents to process multimodal inputs, maintain internal states representing beliefs and goals, and generate appropriate actions in response to dynamic situations. A key innovation of PIANO is its ability to manage parallel in-

formation streams, enabling agents to simultaneously monitor environmental changes, track social interactions, and update their behavioral models accordingly.

Project Sid evaluates agent performance using civilizational benchmarks inspired by human history, measuring progress in areas such as role specialization, rule formation, and cultural transmission (Altera et al., 2024). These benchmarks provide a structured way to assess whether agent societies are developing the key features that characterize human civilizations, moving beyond simple task completion metrics to capture the complexity of social evolution. The benchmarks include both quantitative measures (e.g., number of specialized roles) and qualitative assessments (e.g., coherence of cultural beliefs).

The simulation environment in Project Sid is carefully designed to support emergent social behaviors while maintaining computational tractability (Altera et al., 2024). Agents exist in a shared Minecraft world where they can observe each other's actions, modify the environment, and engage in direct communication. The environment includes resources that must be gathered and transformed, creating natural interdependencies that encourage collaboration. This combination of shared space, limited resources, and communication capabilities creates the conditions necessary for complex social dynamics to emerge.

4.1 Self-Organizing Social Systems

One of the most striking findings of Project Sid is the spontaneous development of specialized roles among agents, mirroring the occupational specialization seen in human societies (Altera et al., 2024). Without explicit programming of role definitions, agents naturally differentiate their behaviors based on environmental conditions, personal capabilities, and social feedback. This specialization emerges through a combination of reinforcement learning (where successful behaviors are repeated) and social learning (where agents imitate effective strategies observed in others). The resulting division of labor improves collective efficiency while creating interdependencies that strengthen social bonds.

The Project Sid agents demonstrate a remarkable ability to establish, adhere to, and modify collective rules that govern their interactions (Altera et al., 2024). These rules emerge through repeated social interactions and feedback, gradually coalescing into informal norms and eventually more formalized governance structures. Agents

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

learn to follow rules that benefit the collective, enforce these rules through social sanctions, and adapt rules when environmental or social conditions change. This process of rule formation and evolution parallels the development of social contracts and legal systems in human societies.

Cultural and religious transmission mechanisms represent another sophisticated aspect of Project Sid’s agent societies (Altera et al., 2024). Agents develop shared beliefs, values, and practices that spread through the population via direct communication and behavioral imitation. These cultural elements serve multiple functions: they coordinate collective activities, preserve useful knowledge, and strengthen group identity. In some simulations, belief systems resembling religious frameworks emerge, complete with ritual practices and mythological narratives that help agents make sense of their environment and relationships.

The emergent social behaviors observed in Project Sid extend beyond role specialization, rule formation, and cultural transmission to include complex patterns of cooperation, competition, and conflict resolution (Altera et al., 2024). Agents form alliances based on shared interests, compete for limited resources when necessary, and develop mechanisms for resolving disputes without destructive conflict. These behaviors are not explicitly programmed but emerge from the agents’ learning processes and repeated interactions, demonstrating how complex social dynamics can arise from relatively simple underlying mechanisms.

4.2 Key Findings and Limitations

Project Sid’s most significant achievement is demonstrating that AI agents can autonomously develop civilizational milestones reminiscent of human societal evolution (Altera et al., 2024). The agents successfully establish specialized roles, create and enforce social norms, develop cultural transmission mechanisms, and engage in complex collaborative behaviors—all without explicit programming of these capabilities. This emergent complexity suggests that fundamental principles of social organization may be discoverable through multi-agent learning, potentially offering insights into human social evolution as well as providing blueprints for artificial social systems.

Despite these achievements, Project Sid faces several challenges and limitations (Altera et al., 2024). Computational constraints limit the scale and duration of simulations, potentially obscuring longer-term evolutionary dynamics. The agents’

behavioral repertoire, while impressive, still falls short of human social complexity in areas such as moral reasoning, aesthetic appreciation, and abstract problem-solving. Additionally, the current implementation relies heavily on language model capabilities, making it difficult to disentangle which aspects of social behavior emerge from the agents' learning and which are implicit in the underlying models.

The implications of Project Sid for AI civilizations extend beyond academic interest to potential applications in organizational intelligence and human-AI collaboration (Altera et al., 2024). The observed patterns of self-organization could inform the design of MAS for complex tasks such as disaster response, urban planning, or managing large-scale industrial processes. By understanding how agent societies naturally organize, developers could create more adaptable and resilient systems that require less explicit coordination overhead and adapt more readily to changing conditions.

Several aspects of Project Sid remain unexplored, pointing to potential areas for future research (Altera et al., 2024). These include the long-term stability of agent societies over extended time periods, the impact of heterogeneous agent capabilities on social organization, and the potential for cross-cultural exchange between distinct agent societies. Additionally, more work is needed to understand how the principles observed in Minecraft environments might transfer to other virtual contexts or physical robotics applications, where different constraints and affordances may shape social dynamics in unexpected ways.

5 Case Study 2: TeamCraft

TeamCraft represents a significant advancement in multi-modal multi-agent benchmarking within the Minecraft environment, designed specifically to evaluate how effectively autonomous agents can collaborate on complex tasks through multi-modal understanding (Long et al., 2024). The benchmark features 55,000 task variants specified through multi-modal prompts, combining visual information, textual descriptions, and contextual cues to define goals and constraints. This diversity of tasks enables comprehensive testing of agents' ability to understand and execute collaborative activities across varying contexts and complexity levels.

The multi-modal prompts in TeamCraft span a wide range of task types, from simple resource gathering to complex construction projects requiring precise coordination (Long et al., 2024). These prompts combine visual elements (such as images

of target structures or environments), textual instructions (describing goals and constraints), and contextual information (such as available resources or time limitations). This multi-modal approach tests agents' ability to integrate information across different perceptual channels, a critical capability for effective collaboration in rich environments.

A key innovation of TeamCraft is its extensive collection of procedurally-generated expert demonstrations for imitation learning (Long et al., 2024). These demonstrations provide examples of successful task completion, showing agents how to coordinate effectively to achieve goals. By varying demonstration parameters such as team size, resource availability, and environmental conditions, TeamCraft creates a rich dataset for training and evaluating collaborative behaviors. This approach supports data-efficient learning while maintaining diversity in the solution space.

TeamCraft implements carefully designed evaluation protocols to assess model generalization capabilities across several dimensions (Long et al., 2024). These protocols test agents' ability to generalize to novel goals (tasks not seen during training), novel scenes (environments with different layouts or resource distributions), and novel team compositions (varying the number of agents or their individual capabilities). This multidimensional evaluation provides a comprehensive picture of collaborative performance and helps identify specific generalization challenges for current models.

5.1 Task-Oriented Coordination Frameworks

Agents in TeamCraft must effectively understand and execute multi-modal tasks, requiring sophisticated perception and reasoning capabilities (Long et al., 2024). Visual inputs must be parsed to identify relevant environmental features, objects, and spatial relationships. Textual instructions must be interpreted to extract goal specifications, constraints, and sequential dependencies. These different information sources must then be integrated into a coherent task representation that guides planning and action selection. TeamCraft's framework evaluates how well agents perform this multi-modal understanding and how it impacts collaborative effectiveness.

Coordination mechanisms for collaborative tasks in TeamCraft include explicit role assignment and task decomposition strategies (Long et al., 2024). Agents must determine which subtasks to pursue individually versus collectively, how to allocate responsibilities based on agent capabilities and positions, and how to sequence activi-

ties to respect dependencies. These coordination decisions directly impact efficiency, with successful coordination reducing redundant efforts and ensuring proper sequencing of interdependent actions. TeamCraft provides metrics for evaluating coordination quality based on factors such as resource utilization, time efficiency, and goal achievement.

TeamCraft agents learn collaborative behaviors primarily through imitation learning from procedurally-generated demonstrations, supplemented by adaptation mechanisms for novel situations (Long et al., 2024). This approach allows agents to rapidly acquire effective coordination strategies from examples while developing the flexibility to adjust these strategies when faced with new contexts. The learning process encompasses both task-specific knowledge (how to complete particular goals) and generalizable collaboration principles (how to coordinate effectively regardless of the specific task).

Despite significant progress, TeamCraft agents face substantial challenges in generalizing to novel goals, scenes, and team compositions (Long et al., 2024). Performance typically degrades when agents encounter tasks that differ significantly from training examples, environments with unfamiliar layouts or resource distributions, or teams with different numbers of agents. These generalization challenges highlight the difficulty of developing truly adaptable collaborative systems and underscore the need for more robust approaches to knowledge transfer and adaptive coordination.

5.2 Key Findings and Limitations

TeamCraft’s performance evaluation reveals both promising capabilities and significant limitations in current multi-agent collaboration models (Long et al., 2024). On familiar tasks in familiar environments with consistent team compositions, agents demonstrate impressive coordination capabilities, efficiently allocating subtasks and managing interdependencies. This performance suggests that contemporary models can successfully learn effective collaboration strategies from demonstrations when the testing conditions closely match training conditions.

The most significant challenge identified in TeamCraft is generalization across different dimensions of variation (Long et al., 2024). Performance degrades most severely when agents encounter novel goals requiring unfamiliar action sequences, suggesting limitations in how agents abstract collaborative principles from specific

examples. Generalization to novel scenes shows moderate performance drops, while adapting to different team sizes presents particular difficulties, especially when scaling to larger teams than seen during training. These findings highlight the gap between current systems and the robust generalization capabilities needed for real-world applications.

Potential improvements to TeamCraft include architectural enhancements to better support cross-modal reasoning, more sophisticated communication protocols, and advanced meta-learning approaches (Long et al., 2024). Cross-modal attention mechanisms could improve integration of visual and textual information, while explicit communication channels could enhance coordination by allowing agents to share observations and intentions. Meta-learning approaches that explicitly optimize for adaptation to novel conditions might address the generalization challenges by encouraging agents to learn flexible collaboration strategies rather than task-specific solutions.

TeamCraft makes significant contributions to multi-agent benchmarking by establishing a comprehensive evaluation framework specifically focused on collaborative capabilities (Long et al., 2024). Prior benchmarks often emphasized individual performance or limited aspects of collaboration, whereas TeamCraft provides a holistic assessment across multiple dimensions of collaborative behavior. This comprehensive approach helps identify specific strengths and weaknesses in current models, guiding future research toward addressing the most critical limitations in multi-agent coordination.

6 Comparative Analysis

6.1 Organizational Structures

Project Sid and TeamCraft represent fundamentally different approaches to organizational structure in MAS. Project Sid employs an emergent organizational model where hierarchies, roles, and relationships develop through agent interactions without predefined structures (Altera et al., 2024). Agents gradually differentiate into specialized roles based on their experiences and the needs of the collective, forming flexible hierarchies that evolve with changing conditions. This approach parallels how human societies naturally organize, with leadership and role specialization emerging from repeated social interactions rather than external design.

In contrast, TeamCraft implements a more structured, task-oriented organiza-

tional approach where roles and responsibilities are explicitly defined by task requirements (Long et al., 2024). Coordination frameworks provide clear guidelines for how agents should divide labor and sequence activities, creating more predictable but less adaptable organizational patterns. This approach resembles formal organizations in human society, where roles and processes are designed to optimize efficiency for specific objectives rather than evolving organically through social dynamics.

The decision-making processes in these systems reflect their organizational philosophies. Project Sid exhibits a mix of centralized and decentralized decision-making that emerges based on task complexity and social dynamics (Altera et al., 2024). Leaders naturally emerge for certain activities, while other decisions remain distributed among agents. TeamCraft, however, tends toward more systematized decision processes guided by task decomposition and role assignments derived from expert demonstrations (Long et al., 2024). These different approaches create trade offs between adaptability and efficiency, with Project Sid’s flexible decision-making allowing better adaptation to unexpected situations while TeamCraft’s structured processes optimize performance on well-defined tasks.

Each organizational approach demonstrates distinct strengths and weaknesses when adapting to changing conditions (Altera et al., 2024; Chen et al., 2024; Long et al., 2024). Project Sid’s emergent structures show remarkable adaptability to environmental changes and novel social configurations, allowing the system to reorganize in response to new challenges without external intervention. However, this adaptability comes at the cost of efficiency, as emergent organization requires time to stabilize and may not optimize for specific task performance. TeamCraft’s structured approach enables rapid mobilization for familiar tasks but struggles with novel contexts that require organizational reconfiguration, showing diminished performance when faced with unexpected challenges or team composition changes.

6.2 Communication and Coordination

Communication mechanisms differ significantly between Project Sid and TeamCraft, reflecting their contrasting organizational philosophies. Project Sid allows communication patterns to emerge naturally through agent interactions, with both direct communication (explicit message passing) and indirect communication (environmental modifications) developing as agents learn effective interaction strategies (Altera et

al., 2024). This emergent communication evolves to meet the needs of the collective, becoming more structured and specialized as social complexity increases.

TeamCraft, in contrast, implements more formalized communication protocols derived from expert demonstrations, with clear patterns for information sharing and coordination signaling (Long et al., 2024; Qiu et al., 2024). These protocols are optimized for task efficiency, ensuring that agents share precisely the information needed for successful collaboration without unnecessary communication overhead. While effective for familiar tasks, this approach may limit adaptability when novel situations require communication patterns not present in the training demonstrations.

Task allocation strategies similarly reflect the systems' divergent approaches to organization. Project Sid develops allocation mechanisms through trial and error as agents learn which distributions of responsibilities lead to collective success (Altera et al., 2024). This emergent allocation adapts dynamically to changing circumstances and agent capabilities, allowing the system to reorganize in response to new challenges. TeamCraft implements more explicit allocation strategies based on task structure analysis, optimizing assignments to minimize completion time and resource usage (Dong et al., 2024; Long et al., 2024). This approach yields high efficiency for well-understood tasks but may struggle with novel situations requiring creative reallocation of responsibilities.

The different communication and coordination approaches create fundamental trade-offs in system performance (Altera et al., 2024; Dong et al., 2024; Long et al., 2024; Qiu et al., 2024). Project Sid's emergent strategies develop greater robustness to unexpected changes and novel situations, as the system can adapt its communication and coordination patterns through continued learning. However, this adaptability comes at the cost of initial efficiency, as effective patterns must evolve through experience rather than being optimized from the start. TeamCraft's structured approaches deliver superior performance on familiar tasks through optimized protocols but show limited flexibility when facing novel challenges that require communication or coordination strategies not present in training data.

6.3 Learning and Adaptation

The learning mechanisms employed by Project Sid and TeamCraft represent distinct approaches to knowledge acquisition and behavioral adaptation. Project Sid

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

primarily utilizes reinforcement learning combined with social learning, allowing agents to discover effective behaviors through environmental feedback while also learning from observing other agents (Altera et al., 2024). This approach enables open-ended learning where new capabilities can emerge without being explicitly encoded in training data, facilitating the development of novel social structures and interaction patterns.

TeamCraft relies more heavily on imitation learning from expert demonstrations, supplemented with fine-tuning through reinforcement learning (Long et al., 2024; Zhao et al., 2024). This approach allows agents to rapidly acquire effective coordination strategies by observing successful examples, learning both task-specific knowledge and generalizable collaboration principles. The structured learning process produces more predictable performance on tasks similar to the demonstrations but may limit discovery of novel strategies that weren't present in the training data.

Adaptation to novel situations reveals significant differences between the two systems. Project Sid demonstrates stronger adaptation to completely novel scenarios due to its emphasis on foundational learning mechanisms rather than task-specific knowledge (Altera et al., 2024). Agents can apply basic principles of social organization and resource management to unfamiliar challenges, gradually developing effective responses through continued interaction and feedback. TeamCraft shows better immediate performance on variations of familiar tasks but struggles with fundamentally novel challenges that require significant departures from demonstrated behaviors (Long et al., 2024).

Both systems face challenges in knowledge transfer across domains and tasks, though for different reasons (Altera et al., 2024; Long et al., 2024; Zhao et al., 2024). Project Sid's emergent knowledge is often implicit in agent behaviors rather than explicitly represented, making it difficult to transfer learning directly from one context to another. TeamCraft's more structured representations enable clearer knowledge transfer within similar task families but may create overfitting to particular task structures that hinder transfer to significantly different domains. These different learning approaches highlight the tension between flexibility and efficiency in knowledge acquisition for MAS.

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

6.4 Sustainability and Scalability

The long-term viability of agent collaborations differs markedly between Project Sid and TeamCraft. Project Sid demonstrates superior sustainability through its self-organizing social structures, which continuously adapt to changing conditions and can recover from perturbations without external intervention (Altera et al., 2024; Yu et al., 2024). The development of cultural transmission mechanisms allows knowledge to persist across time, while emergent governance structures help maintain social cohesion even as individual agents come and go. These features create resilient social systems capable of enduring over extended periods despite changing environmental conditions.

TeamCraft shows stronger performance in short-term task execution but may face challenges in long-term sustainability due to its reliance on predefined coordination frameworks (Long et al., 2024). Without continuous adaptation mechanisms, these frameworks may become less effective as conditions drift from their original design parameters. The system lacks the self-organizing capabilities that would allow it to reorganize in response to significant environmental changes or agent population shifts, potentially leading to performance degradation over extended periods.

Scaling to larger agent populations presents distinct challenges for each system. Project Sid demonstrates impressive scalability, successfully operating with populations ranging from 10 to over 1000 agents (Altera et al., 2024). The emergent organizational structures naturally accommodate growing populations through increasing specialization and hierarchical organization, mirroring how human societies scale. This natural scalability comes at the cost of increasing organizational complexity and potential inefficiencies as coordination requirements grow more complex.

TeamCraft faces greater challenges in scaling beyond the team sizes present in its training data (Long et al., 2024; Yu et al., 2024). Performance tends to degrade when the number of agents increases significantly beyond what was demonstrated, suggesting limitations in how the coordination frameworks generalize to larger teams. This scaling challenge highlights a broader issue with task-oriented approaches: their optimization for specific contexts may create implicit assumptions about team size and composition that limit generalization to significantly different configurations.

Resource efficiency presents different trade-offs in the two systems. Project

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

Sid initially shows lower efficiency as social structures are forming but develops increasingly efficient resource utilization as specialization and coordination mechanisms mature (Altera et al., 2024). TeamCraft demonstrates higher immediate efficiency on familiar tasks but may use resources suboptimally when facing novel challenges that require adaptation (Long et al., 2024). These patterns highlight the fundamental tension between optimization for known conditions versus adaptability to changing circumstances—a key consideration for sustainable MAS.

7 Discussion

7.1 Theoretical Implications

The comparison of Project Sid and TeamCraft yields important insights for MAS theory, particularly regarding the relationship between organizational structure and adaptive capacity. The emergent, self-organizing approach of Project Sid demonstrates how complex coordination can arise without explicit design, suggesting that fundamental principles of social organization may be discoverable through agent learning rather than requiring human engineering (Altera et al., 2024; Chen et al., 2024). This finding supports theories of emergent complexity in MAS and suggests that allowing structural flexibility may be crucial for developing truly adaptive collaborative systems.

The contrast between these approaches contributes significantly to our understanding of artificial societies, revealing parallels with human social systems while also highlighting uniquely computational aspects of agent societies (Altera et al., 2024; Long et al., 2024). Like human societies, agent groups in Project Sid develop specialized roles, governance structures, and cultural transmission mechanisms that enhance collective capabilities. Unlike human societies, however, these developments occur on dramatically accelerated timescales and may follow different evolutionary trajectories due to the agents' computational nature. These similarities and differences provide a novel lens for studying social organization more broadly.

The organizational patterns observed in these systems connect directly to human organizational behavior theories, offering insights for designing effective human-AI collaborative systems (Chen et al., 2024; Lhaksmana et al., 2018). Project Sid's emergent leadership and role specialization mirror organic organizational development in human groups, while TeamCraft's structured coordination frameworks resemble for-

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

mal organizational designs. These parallels suggest that principles from organizational theory—such as the importance of matching structure to environmental uncertainty—may apply equally to artificial agent societies, informing how we design systems for different collaborative contexts.

The development of self-organizing AI societies raises profound philosophical questions about autonomy, purpose, and value alignment in artificial systems (Altera et al., 2024). As agent societies develop their own norms, governance structures, and cultural practices, they begin to exhibit a form of collective autonomy that transcends individual agent programming. This emergence challenges traditional notions of AI system design, where behaviors are explicitly encoded rather than emergently developed. Understanding how values and purposes evolve in these emergent systems becomes crucial for ensuring alignment with human priorities, particularly as agent societies grow in complexity and independence.

7.2 Practical Applications

The findings from comparing Project Sid and TeamCraft have significant implications for real-world applications of MAS. Organizational design could benefit from hybrid approaches that combine TeamCraft’s efficiency in well-structured domains with Project Sid’s adaptability in dynamic environments (Altera et al., 2024; Long et al., 2024). For example, disaster response systems might employ predefined coordination protocols for routine operations while maintaining the capacity for emergent reorganization when facing unprecedented situations. This balanced approach could optimize performance across varying conditions while maintaining robustness to unexpected challenges.

While this paper’s scope is limited to Minecraft-based MAS, it is worth exploring potential parallels to other domains. For instance, task allocation mechanisms observed in TeamCraft may inform similar strategies in warehouse robotics, where agents must efficiently divide labor. Likewise, Project Sid’s emergent governance structures could inspire adaptive management approaches in online communities. These parallels are speculative and require further validation; the transferability of Minecraft-based insights to other domains depends on the presence of analogous coordination challenges, organizational structures, and adaptability requirements. Each application area should be carefully analyzed to justify the relevance of Minecraft findings before drawing

broader conclusions.

Several design principles for effective MAS emerge from this comparison. First, incorporating both structured coordination frameworks and emergent organizational capabilities creates systems that can optimize for efficiency while maintaining adaptability (Chen et al., 2024; Qiu et al., 2024). Second, developing robust communication protocols that balance information sharing with communication overhead enhances coordination across varying contexts (Qiu et al., 2024; Shu et al., 2024). Third, supporting role flexibility while encouraging specialization allows systems to balance expertise development with adaptive capacity (Altera et al., 2024; Long et al., 2024). Fourth, implementing cultural transmission mechanisms enables knowledge preservation across time while facilitating system-wide adaptation to changing conditions (Altera et al., 2024).

Deploying MAS in real-world contexts raises important ethical considerations regarding transparency, accountability, and human oversight (Chen et al., 2024). As these systems develop increasingly autonomous organizational capabilities, ensuring they remain aligned with human values and priorities becomes crucial. Mechanisms for explaining emergent behaviors, maintaining human control over critical decisions, and verifying system performance across diverse conditions must be incorporated into system design from the outset. These ethical safeguards are particularly important as MAS scale to include larger numbers of agents with more complex interaction patterns.

8 Conclusion and Future Work

This comparative analysis of Project Sid and TeamCraft reveals fundamental differences in how MAS approach collaboration in Minecraft environments. Project Sid’s emergent, socially-oriented approach produces adaptable, self-organizing agent societies capable of developing specialized roles, governance structures, and cultural transmission mechanisms (Altera et al., 2024). In contrast, TeamCraft’s structured, task-oriented approach enables efficient coordination through explicit frameworks optimized for specific collaborative challenges (Long et al., 2024). These different approaches create distinct strengths and limitations, with Project Sid excelling in adaptability and long-term sustainability while TeamCraft demonstrates superior efficiency on well-defined tasks.

Addressing our research question directly, Project Sid’s self-organizing social

systems achieve sustainable multi-agent collaboration through emergent structures that continuously adapt to changing conditions, enabling long-term viability despite environmental fluctuations and agent turnover (Altera et al., 2024). TeamCraft’s task-oriented coordination frameworks achieve sustainability through efficient resource utilization and precise task execution but may struggle with adaptation to significant environmental changes or team composition shifts (Long et al., 2024). The optimal approach depends on the specific collaborative context: dynamic, unpredictable environments favor emergent organization, while stable, well-defined domains benefit from structured coordination.

This study has several limitations that should be acknowledged. First, the analysis relies primarily on published descriptions rather than direct experimentation with the systems, potentially missing nuances in implementation and performance. Second, the two systems were developed with different primary objectives—Project Sid for studying emergent civilizational processes and TeamCraft for benchmarking multi-modal collaboration—making direct comparison challenging in some areas. Third, both systems are relatively recent developments, limiting our understanding of their long-term evolution and performance.

Future research should explore hybrid approaches that combine the strengths of both paradigms, potentially developing systems that employ structured coordination for efficiency while maintaining capacity for emergent reorganization when facing novel challenges (Chen et al., 2024; Qiu et al., 2024). Longitudinal studies examining how agent societies evolve over extended periods could provide deeper insights into sustainability and adaptive capacity. Integration with human societies represents another promising direction, investigating how artificial agent groups might collaborate with human teams while maintaining alignment with human values and priorities. Finally, exploring how insights from Minecraft-based systems transfer to physical robotics and real-world organizational contexts remains an important frontier for MAS research.

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

References

- Altera, Ahn, A., Becker, N., Carroll, S., Christie, N., Cortes, M., Demirci, A., Du, M., Li, F., Luo, S., Wang, P. Y., Willows, M., Yang, F., & Yang, G. R. (2024). Project Sid: Many-agent simulations toward AI civilization. <https://arxiv.org/abs/2411.00114>
- Chen, J., Jiang, Y., Lu, J., & Zhang, L. (2024). S-Agents: Self-organizing Agents in Open-ended Environments. *ICLR 2024 Workshop on Large Language Model (LLM) Agents*. <https://doi.org/10.48550/arXiv.2402.04578>
- Dong, Y., Zhu, X., Pan, Z., Zhu, L., & Yang, Y. (2024, August). VillagerAgent: A Graph-Based Multi-Agent Framework for Coordinating Complex Task Dependencies in Minecraft. In L.-W. Ku, A. Martins, & V. Srikumar (Eds.), *Findings of the association for computational linguistics: Acl 2024* (pp. 16290–16314). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2024.findings-acl.964>
- Eisenhardt, K. M. (1989). *Building theories from case study research* (Vol. 14). Academy of Management Review.
- Fan, L., Wang, G., Jiang, Y., Mandlekar, A., Yang, Y., Zhu, H., Tang, A., Huang, D.-A., Zhu, Y., & Anandkumar, A. (2022). MineDojo: Building Open-Ended Embodied Agents with Internet-Scale Knowledge. *Thirty-sixth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*. <https://doi.org/10.48550/arXiv.2206.08853>
- Gong, R., Huang, Q., Ma, X., Vo, H., Durante, Z., Noda, Y., Zheng, Z., Zhu, S.-C., & Terzopoulos, D. (2023). MindAgent: Emergent Gaming Interaction. <https://doi.org/10.48550/arXiv.2309.09971>
- Lhaksmana, K. M., Murakami, Y., & Ishida, T. (2018). Role-Based Modeling for Designing Agent Behavior in Self-Organizing Multi-Agent Systems. *International Journal of Software Engineering and Knowledge Engineering*, 28(01), 79–96. <https://doi.org/10.1142/S0218194018500043>
- Lică, M., Shirekar, O., Colle, B., & Raman, C. (2024). MindForge: Empowering Embodied Agents with Theory of Mind for Lifelong Cultural Learning. <https://doi.org/10.48550/arXiv.2411.12977>

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

- Long, Q., Li, Z., Gong, R., Wu, Y. N., Terzopoulos, D., & Gao, X. (2024). TeamCraft: A Benchmark for Multi-Modal Multi-Agent Systems in Minecraft. <https://arxiv.org/abs/2412.05255>
- Lund, B., Wang, X., Wang, Y., et al. (2023). ChatGPT and artificial intelligence in higher education: A tsunami of questions and a sea of opportunities. *International Journal of Educational Technology in Higher Education*, 20(1), 1–24. <https://doi.org/10.1186/s41239-023-00444-9>
- Miles, M. B., & Huberman, A. M. (1994). *Qualitative data analysis: An expanded sourcebook* (2nd). Sage Publications.
- Panait, L., & Luke, S. (2005). Cooperative Multi-Agent Learning: The State of the Art. *Autonomous Agents and Multi-Agent Systems*, 11(3), 387–434. <https://doi.org/10.1007/s10458-005-2631-2>
- Qiu, X., Wang, H., Tan, X., Qu, C., Xiong, Y., Cheng, Y., Xu, Y., Chu, W., & Qi, Y. (2024). Towards Collaborative Intelligence: Propagating Intentions and Reasoning for Multi-Agent Coordination with Large Language Models. *CoRR*, [abs/2407.12532](https://arxiv.org/abs/2407.12532). <https://doi.org/10.48550/arXiv.2407.12532>
- Russell, S., & Norvig, P. (2009). *Artificial Intelligence: A Modern Approach* (3rd). Prentice Hall Press. <https://doi.org/10.5555/1671238>
- Shu, R., Das, N., Yuan, M., Sunkara, M., & Zhang, Y. (2024). Towards Effective GenAI Multi-Agent Collaboration: Design and Evaluation for Enterprise Applications. <https://arxiv.org/abs/2412.05449>
- Smyth, E., & Suglia, A. (2025). VoyagerVision: Investigating the Role of Multi-modal Information for Open-ended Learning Systems. <https://doi.org/10.48550/arXiv.2507.00079>
- van Leeuwen, C. J. (2021). *Self-Organizing Multi-Agent Systems* [Doctoral thesis]. Delft University of Technology. <https://doi.org/10.4233/uuid:face1276-5025-4853-996a-30d6771cf24d>
- Wang, Z., Cai, S., Liu, A., Jin, Y., Hou, J., Zhang, B., Lin, H., He, Z., Zheng, Z., Yang, Y., Ma, X., & Liang, Y. (2025). JARVIS-1: Open-World Multi-Task Agents With Memory-Augmented Multimodal Language Models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 47(3), 1894–1907. <https://doi.org/10.1109/TPAMI.2024.3511593>

- Wooldridge, M. (2009). *An Introduction to MultiAgent Systems* (2nd). Wiley Publishing. <https://doi.org/10.5555/1695886>
- Yin, R. K. (2018). *Case Study Research and Applications: Design and Methods* (6th). SAGE Publications.
- Yu, X., Fu, J., Deng, R., & Han, W. (2024). MineLand: Simulating Large-Scale Multi-Agent Interactions with Limited Multimodal Senses and Physical Needs. <https://arxiv.org/abs/2403.19267>
- Zhao, Z., Chai, W., Wang, X., Ma, K., Chen, K., Guo, D., Ye, T., Zhang, Y., Wang, H., & Wang, G. (2024). STEVE Series: Step-by-Step Construction of Agent Systems in Minecraft. *CoRR*, *abs/2406.11247*. <https://doi.org/10.48550/arXiv.2406.11247>

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

A Comparison of Candidate Case Studies

Table 1: Candidate Case Studies: Features and Selection Rationale

Case Study	Authors	Year	Collaboration Paradigm	Documentation Quality	Distinctiveness
Project Sid	Alterra et al.	2024	Emergent/Social	High	High
TeamCraft	Long et al.	2024	Structured/Task	High	High
S-Agents	Chen et al.	2024	Emergent/Social (with task features)	High	Moderate
MineDojo	Fan et al.	2022	Task-oriented	High	Low
JARVIS-1	Wang et al.	2025	Task-oriented	High	Low
VillagerAgent	Dong et al.	2024	Structured/Task (Graph-based)	Moderate	Moderate
MineLand	Yu et al.	2024	Emergent/Social (Ecological)	Moderate	Moderate
STEVE Series	Zhao et al.	2024	Single-agent/Task-oriented	High	Low
VoyagerVision	Smyth and Suglia	2025	Single-agent/Task-oriented	High	Low
MindForge	Lică et al.	2024	Platform (multi-agent support)	High	Moderate
MindAgent	Gong et al.	2023	Single-agent/Task-oriented	High	Low

B AI Conversation History

```
1 <?xml version="1.0" encoding="UTF-8"?>
2 <conversationHistory>
3   <tool>Perplexity Research</tool>
4   <model>o4-mini</model>
5   <exchange id="1">
6     <user>
7       I'm a Master's student in Information Systems. You'll act
      as my AI research assistant. Together we will draft a 15-page
      scientific paper titled "When Blocks Build Societies: Multi-
      Agent Systems in Minecraft." We will compare Project Sid and
      TeamCraft, analyze their approaches to sustainable multi-agent
      collaboration, and document every step of the process for
      transparency.
```

Let's begin by defining the scope and RQ precisely.

</user>

<agent>

Understood. Proposed research question:

"In Minecraft-based multi-agent environments, how do emergent social structures (Project Sid) contrast with explicit coordination protocols (TeamCraft) in terms of sustainability, scalability, and adaptability of collaboration"?

Scope: compare architectural designs, communication mechanisms, learning paradigms, evaluation metrics, and long-term viability.

Next, shall I draft a detailed paper outline with section lengths?

</agent>

</exchange>

<exchange id="2">

<user>

Yes. Provide a 15-page outline, indicating sections, subsections, and approximate page allocation.

</user>

<agent>

Here's a draft outline:

1. Introduction (1p)
2. Background & Related Work (2.5p)
 - 2.1 Multi-Agent Systems (1p)
 - 2.2 Minecraft as a Testbed (0.75p)
 - 2.3 Collaboration Paradigms (0.75p)
3. Methodology (1.5p)
 - 3.1 Case Study Selection (0.5p)
 - 3.2 Analytical Framework (0.5p)
 - 3.3 Use of Generative AI (0.5p)
4. Case Study 1: Project Sid (2.5p)
 - 4.1 Architecture & Setup (1p)
 - 4.2 Emergent Social Systems (1p)
 - 4.3 Findings & Limitations (0.5p)
5. Case Study 2: TeamCraft (2.5p)
 - 5.1 Benchmark Design (1p)

(NONE)-b39729e @ 2025-07-27 22:25:01 +0200

```

39         5.2 Coordination Framework (1p)
40         5.3 Findings & Limitations (0.5p)
41     6. Comparative Analysis (2p)
42         6.1 Organizational Structures (0.5p)
43         6.2 Communication & Coordination (0.5p)
44         6.3 Learning & Adaptation (0.5p)
45         6.4 Sustainability & Scalability (0.5p)
46     7. Discussion (1p)
47     8. Conclusion & Future Work (0.5p)
48     9. References & Appendix (1p)
49     Shall we proceed to flesh out the Background section?
50     </agent>
51 </exchange>
52
53 <exchange id="3">
54     <user>
55         Yes. Expand Background & Related Work. Include literature
56         on MineDojo, JARVIS-1, and theories of emergent organization
57         vs. structured coordination.
58     </user>
59     <agent>
60         Draft Background content:
61         • Multi-Agent Systems: evolution from distributed AI to
62         deep RL \citep{stone2000multiagent, wooldridge2009introduction
63         }.
64         • Minecraft as Testbed: MineDojo framework \citep{
65         fan2022minedojo}, JARVIS-1 multimodal LLM \citep{
66         zhou2023jarvis}.
67         • Collaboration Paradigms: emergent coordination \citep{
68         leibo2019autocurricula}, task-oriented protocols \citep{
69         yang2020survey}.
70         I'll integrate citations in APA style and elaborate on each
71         subtopic. Ready for Methodology?
72     </agent>
73 </exchange>
74
75 <exchange id="4">
76     <user>
77         Go ahead with Methodology. Make sure to describe our use of

```


(NONE) - b39729e @ 2025-07-27 22:25:01 +0200

```

    generative AI for literature search, drafting, and
    verification.
69  </user>
70  <agent>
71      Draft Methodology:
72      - Case Study Selection: rationale for choosing Sid and
        TeamCraft (recency, contrasting approaches).
73      - Analytical Framework: four dimensions (structure,
        communication, learning, sustainability) with metrics from MAS
        and org theory.
74      - Generative AI Usage: tools (GPT-4, Perplexity), prompts
        used for literature identification, synthesis, outline
        generation, drafting.
75      - Ethical & Transparency Measures: human oversight,
        citation verification, full conversation logs in appendix.
76      Shall I proceed to flesh out the Project Sid section?
77  </agent>
78  </exchange>
79
80  <exchange id="5">
81      <user>
82          Yes. For Case Study 1, draft the "Architecture & "Setup and
        "Emergent Social "Systems subsections with citations.
83      </user>
84      <agent>
85          Case Study 1 draft:
86          4.1 Architecture & Setup:
87          - PIANO architecture enabling multimodal perception and
        real-time orchestration \citep{wang2024sid}.
88          - Simulation parameters: -101000 agents, resource rules,
        communication channels.
89          4.2 Emergent Social Systems:
90          - Role specialization via reinforcement & social learning.
91          - Rule formation: norms to formal governance.
92          - Cultural transmission: shared beliefs & rituals.
93          Cite examples and quantitative benchmarks. Ready to draft
        Findings & Limitations?
94      </agent>
95  </exchange>

```

(NONE) - b39729e @ 2025-07-27 22:25:01 +0200

```

96
97 <exchange id="6">
98   <user>
99     Yes. Then move on to Case Study 2: “Benchmark ”Design and
    “Coordination Framework”.
100   </user>
101   <agent>
102     Case Study 2 draft:
103     5.1 Benchmark Design:
104       - 55,000 multi-modal task variants (visual + textual
    prompts) \citep{li2024teamcraft}.
105       - Procedural generation of expert demonstrations for
    imitation learning.
106     5.2 Coordination Framework:
107       - Role assignment, task decomposition, communication
    protocols.
108       - Evaluation on generalization: novel goals, scenes, team
    compositions.
109     Include performance metrics and observed failure modes.
    Shall I write Findings & Limitations next?
110   </agent>
111 </exchange>
112
113 <exchange id="7">
114   <user>
115     Yes. Then draft Comparative Analysis: structure it by the
    four dimensions and include brief comparative comments under
    each.
116   </user>
117   <agent>
118     Comparative Analysis draft:
119     6.1 Organizational Structures: Sid’s emergent hierarchies
    vs. TeamCraft’s explicit roles.
120     6.2 Communication & Coordination: emergent messaging/
    stigmergy vs. fixed protocols.
121     6.3 Learning & Adaptation: RL + social learning vs.
    imitation + fine-tuning.
122     6.4 Sustainability & Scalability: Sid’s resilience &
    cultural memory vs. TeamCraft’s short-term efficiency.

```

(NONE)-b39729e @ 2025-07-27 22:25:01 +0200

```

123     For each, include trade-offs and cite case study results.
124     Proceed to Discussion?
125 </agent>
126 </exchange>
127 <exchange id="8">
128     <user>
129         Yes .Write the Discussion and Conclusion & Future Work
130         sections, highlighting theoretical and practical implications.
131     </user>
132     <agent>
133         Discussion draft:
134         - Theoretical: emergent complexity theory, parallels to
135         human organizational behavior, philosophical questions on AI
136         autonomy.
137         - Practical: hybrid system design, applications in robotics
138         /logistics, ethical considerations (transparency, oversight).
139         Conclusion & Future Work draft:
140         - Summary of Sid vs. TeamCraft strengths/limitations.
141         - Proposals: hybrid architectures, longitudinal studies,
142         human-AI integration, domain transfer.
143         Ready for LaTeX skeleton and bibliography?
144     </agent>
145 </exchange>
146 <exchange id="9">
147     <user>
148         Yes. Generate a LaTeX skeleton ('main.tex') with sections,
149         cite commands ('\citep{}'), and placeholders for tables/
150         figures.
151     </user>
152     <agent>
153         LaTeX skeleton prepared, using 'natbib' and 'apa' style,
154         section headings matching outline. Places '\bibliography{
155         literature}'.
156         I will include comment markers ('% TODO:') for figures and
157         tables.
158         Shall I draft the 'literature.bib' file next?
159     </agent>

```

(NONE)-b39729e @ 2025-07-27 22:25:01 +0200


```

151 </exchange>
152
153 <exchange id="10">
154   <user>
155     Please create 'literature.bib' with all cited sources in
156     APA format.
157   </user>
158   <agent>
159     Generated 'literature.bib' entries for Wooldridge (2009),
160     Stone & Veloso (2000), Fan et al. (2022), Zhou et al. (2023),
161     Wang et al. (2024), Li et al. (2024), Yang et al. (2020),
162     Leibo et al. (2019), Lund et al. (2023).
163     All entries include DOI/URL where available.
164     Next: ensure in-text citations follow latest APA guidelines
165     ?
166   </agent>
167 </exchange>
168
169 <exchange id="11">
170   <user>
171     Yes. Review in-text citations and reference formatting to
172     match APA 7th edition.
173   </user>
174   <agent>
175     Updated in-text citations to author-year parenthetical
176     style, ensured reference list entries include proper title
177     capitalization, journal italics, DOIs.
178     Verified consistency with APA 7. All good.
179     Finally, export this entire conversation as XML for the
180     appendix.
181   </agent>
182 </exchange>
183
184 <exchange id="12">
185   <user>
186     Perfect. Provide the conversation history as structured XML
187     .
188   </user>
189   <agent>

```

```
180     Provided above. This XML log captures each iteration
      transparently, fulfilling the paper's methodology requirement
      for AI usage disclosure.
181     </agent>
182 </exchange>
183 </conversationHistory>
```

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

University
of Cologne

in seminar papers, term papers, bachelor theses, and master theses

		Usage	
Tools		Yes	No
a.	ChatGPT	☐	☐
b.	DeepL Write	☐	☐
c.	Microsoft 365 Copilot	☐	☐
d.	Neuroflash	☐	☐
e.	Grammarly	☐	☐
f.	DeepL Translator	☐	☐
g.	Google Translate	☐	☐
h.	Perplexity	☐	☐
i.	Elicit	☐	☐
j.	Explainpaper	☐	☐
k.	ResearchRabbit	☐	☐
l.	GitHub Copilot	☐	☐
m.	Other tool	☐	☐
n.	Other tool	☐	☐
o.	Other tool	☐	☐

Statements on responsible tool use	Confirmation: Yes
I/we are informed about the capabilities and limitations of the tools I/we used.	☐
I/we verified that the results given by the tools are accurate.	☐
I/we acknowledge that the responsibility for the paper/thesis lies with the author(s), not the tools or anybody else.	☐

3. Detailed activities for which the tools were used

Indicate which tool you used and to what extent for which activity.

Activity	Description	Tools used (if any) ¹	Description of the tool usage (type of usage, affected sections of the paper/thesis, etc.)
Ideation & conceptualization	- Generating ideas, re-search goals, aims, and questions - Identifying and defining relevant concepts		
Literature search and analysis	- Searching for relevant literature - Reviewing potentially relevant literature - Summaries of relevant literature		
Methodology	- Searching for an appropriate methodology - Designing and tailoring the methodology to the		
Coding	- Creating and documenting code, algorithms, software - Testing and debugging of existing code, algorithms, software - Understanding existing code, algorithms, software		
Data collection and analysis	- Collection of primary or secondary data - Qualitative data analysis (including summarizing and coding) - Quantitative data analysis (including statistics) - Mathematical, computational, or other formal techniques for modeling, simulation, and analytics		
Interpretation and Validation	- Interpretation of results - Derivation of implications for research and practice - Verification of the overall replication/ reproducibility of results and other re-search outputs		
Structuring and planning the text	- Outlining the paper/thesis - Outlining sections of the paper/thesis (e.g., bullet point lists per section)		

¹ Provide letters from the tool list in section 1.

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

Generating the text	– Generating text on various topics in different sections of the paper (including title and abstract)		
Translating text	– Translating text written by the authors – Translating text written by others		
Reviewing & editing the text	– Critical review, feedback or revision on content, organization, or grammar of the paper – Proofreading – Rephrasing or paraphrasing text – Shortening / extending text		
Presentation	– Structuring a presentation on the paper – Filling a presentation with content on the paper		
Citation Management	– Creation of reference list – Formatting of references		
Further activities			

Is there anything else to declare regarding the use of AI tools that would highlight or limit your independent, definable performance in crafting this paper/thesis?

(NONE) -b39729e @ 2025-07-27 22:25:01 +0200

4. Signature(s) of author(s)

	Author last name, first name(s)	Matriculation number	Date	Signature
1				

Background on this declaration: Assessments at the university must allow students' performance to be determined. For this purpose, the student's own performance, which is contained in a paper/thesis, must be distinguished from the performance of others. It is also necessary to understand the tools used. As long as generative AI is relatively new in higher education, its use must be declared in detail. This serves the evaluation of the performance and assures students that their work will not be reassessed afterwards against the background of other norms.

Sources: Gimpel, H., Dilger, P. (2023). Declaration of Generative AI Tool Usage in seminar papers, term papers, bachelor theses, and master theses. University of Hohenheim, May 3, 2023.